Automatic Protoboard Layout

by

Michael Mekonnen

B.S. EECS, Massachusetts Institute of Technology (2013)
B.S. Mathematics, Massachusetts Institute of Technology (2013)

Submitted to the Department of Electrical Engineering and Computer
Science
in partial fulfillment of the requirements for the degree of

Masters of Engineering in Electrical Engineering and Computer
Science

at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

February 2014

Author . . . .
Department of Electrical Engineering and Computer Science
November 4, 2013

Certified by. . . . .
Dennis M. Freeman
Professor
Thesis Supervisor

Certified by. . . . .
Adam J. Hartz
Lecturer
Thesis Supervisor

Accepted by . . . .
Professor Albert R. Meyer
Chairman, Department Committee on Graduate Theses

Automatic Protoboard Layout

by

Michael Mekonnen

Submitted to the Department of Electrical Engineering and Computer Science
on November 4, 2013, in partial fulfillment of the
requirements for the degree of
Masters of Engineering in Electrical Engineering and Computer Science

Abstract

Abstract.

Thesis Supervisor: Dennis M. Freeman
Title: *Professor*

Thesis Supervisor: Adam J. Hartz
Title: *Lecturer*

Acknowledgments

Acknowledgments.

Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Problem Statement

What problem are we solving?

1.2 Motivation

Why is this an interesting problem?

1.3 Goal

Precisely state the goal of this project. In particular, explain that we ultimately want to make a teaching tool for 6.01.

1.4 Outline

How is the Thesis organized?

Chapter 2

Background

In this section we will discuss essential background information to this project. First we discuss the specific terminology used in this paper. Next we discuss previous work done that relates to this project.

2.1 Technical Background

As already mentioned, this project aims to produces a new teaching tool for the introductory course 6.01. We will now discuss the scope of circuits in 6.01.

2.1.1 What are our circuit components?

The rudimentary circuit components that students work with are resistors, operational amplifiers (op amps), and potentiometers (pots). Students start out by building very simple circuits, and then go on to building more complicated circuits with time. The simplest circuits that students build aim to control lego motors in a particular way. In constructing these circuits, students use 6-pin *motor connectors* to connect their circuits to lego motors. The more complicated circuits students build interact with robots that were built specifically for the purposes of 6.01. One of the 6.01 robots is displayed in Figure 2-1. Students use 8-pin *robot connectors* to connect their circuits to robots. The robots can be equipped with heads that have vision

Figure 2-1: 6.01 robot.

capabilities. Each head has a rod attached to a potentiometer. Also attached to the rod are a lego motor and a plate containing two photosensors positioned at a $90°$ angle from each other. The photosensors are used to serve as eyes for the robot. This setup allows us to turn the head by controlling the motor and inquire the current position of the head by probing the pot. Figure 2-1 displays a robot with a head. Students use 8-pin *head connectors* to connect their circuits to robot heads.

All together, our components are resistors, op amps, pots, motor connectors, robot connectors, and head connectors.

2.1.2 What is a circuit schematic?

Throughout this paper, the term *circuit schematic* will refer to a drawing or a sketch of a circuit containing its components and all the interconnections between the components drawn as wires. This is what one would sketch on a piece of paper in the process of designing a circuit. Figure 2-2 presents an example of a circuit schematic.

2.1.3 What is a protoboard?

Protoboards are constructs that make it easy to quickly build and test small circuits. They present a 2-dimensional array of cleverly interconnected dots in which circuit pieces and wires can be inserted. Figure 2-3 presents an example of a physical protoboard. In the orientation depicted in Figure 2-3, a protoboard has 4 groups of rows:

Figure 2-2: Sample circuit schematic.

Figure 2-3: Physical protoboard.

the first 2 rows, the next 5 rows, the next 5 rows, and finally the last 2 rows. In the first and last groups, the dots on the protoboard are interconnected horizontally. In the middle two groups, the dots on the protoboard are interconnected vertically. This interconnection scheme is depicted in Figure 2-3.

2.1.4 What is a protoboard layout of a circuit schematic?

The protoboard layout of a given schematic is the placement of circuit pieces and wires on a protoboard that corresponds to the schematic. This is done by placing the appropriate pieces on the protoboard and then appropriately interconnecting them with wires as prescribed by the schematic. As an example, Figure 2-4 presents the protoboard layout corresponding to the example schematic shown in Figure 2-2.

For each of the circuit components we are interested in, there is a corresponding

Figure 2-4: Protoboard layout for the schematic in Figure 2-2.

circuit piece that may be inserted into the protoboard. The one exception is that op amps come in pairs. That is, each op amp circuit piece that is inserted in the protoboard actually contains two op amp components within it. This raises an important design question when we layout a schematic: what is the best way to group together the op amps in a schematic to result in the "best" layout? In answering this question, the designer must have some criteria for what makes a layout "good." While there are no conclusive answers for this question, general rules of thumb are (in no particular order):

- The layout should have no crossing wires.

- The layout should not have any wires that cross circuit pieces.

- The layout should only have horizontal and vertical wires.

- The layout should have as few wires as possible.

- The total length of wires in a layout should be as small as possible.

Given the background information discussed thus far, the goal of our project is generating a "good" protoboard layout from circuit schematics automatically.

2.2 Previous Work

Here we will discuss previous work that has been done relating to this project. First, as our project aims to augment the quality of 6.01, we look at the current infrastructure available for students. Next, we look at what work has been done relating to layout in general.

2.2.1 CMax

In a typical circuits lab in 6.01, students first design a circuit by drawing a schematic of the circuit on paper and discussing their design with a staff member. After they iteratively amend their design and are happy with it, they build the circuit on a simulation tool called Circuits Maximus (CMax). With this tool, students can layout their circuits on a simulated protoboard as if they were laying it out on a physical one. CMax allows students to test the circuit to make sure that it behaves as desired. Circuit layout is much easier on CMax than on a physical protoboard. Hence, CMax provides a very fast and safe way of debugging circuit layouts. Once the students are satisfied with their observations from CMax, they build their circuits on physical protoboards and carryout the appropriate experiments.

CMax has been a fantastic resource for 6.01 students. Its introduction has made learning circuits significantly easier for many students, especially those that have little or no prior experience with circuits. In addition to making the lab exercises much more manageable, it provides students with a very handy way to build, analyze, and experiment with circuits at their own leisure outside of lab.

While CMax is a fantastic tool, we can imagine a tool that can be even more useful. The most instructive part of the labs that students do in the circuits module of 6.01 is really designing the circuits in the first place, which they currently do by drawing schematic diagrams on paper. Once they are happy with their schematic diagrams, they proceed to laying out the corresponding circuits with CMax. The process of laying out a schematic does not really have very much instructive substance. This process is essentially solving a puzzle, and has almost nothing to do with the subject

14

matter – designing circuits. In fact, when the circuits get complicated and involve many pieces, translating a schematic diagram into a protoboard layout gets to be quite challenging and time-consuming. In these situations, students often end up with convoluted and unpleasant layouts that are very difficult to debug in the likely case of the circuit not behaving as expected.

In the best case scenario, students should not have to produce protoboard layouts for their schematic diagrams. Indeed they should work out the right schematic diagram of the circuit of interest, but the layout generation should not be part of the learning process. This project aims to let students draw and analyze schematic drawings of circuits and produce the corresponding protoboard layouts automatically. Given the protoboard layouts output by this tool, students can proceed to building the circuits on physical protoboards and carrying out the appropriate experiments.

With this tool, a typical circuits lab would go as follows. First, as before, the students draw schematic diagrams of their circuits on paper. Once they have schematic drawings they are happy with, they can directly draw their schematic drawings on the simulation tool. In fact, students my go straight to building the schematic drawings on the simulation tool, bypassing the experimentation on paper. Once they have a schematic drawn, they can analyze it with the tool, discuss it with staff members, and amend it easily and quickly with the user-friendly graphical user interface of the simulation tool. When they are satisfied with the behaviors of their schematic circuit, they can produce the corresponding protoboard layout simply at the click of a button – this would be the most important advantage of this tool. They can then build the layout on a physical protoboard and carryout experiments with it.

2.2.2 Current work in automatic layout

In my explorations, I was not able to find any tools that completely automatically convert circuit schematics into protoboard layouts. However, there do exists tools that perform partially- or fully-automatic Printed Circuit Board (PCB) layout. To my findings, most of these tools do not publish their algorithms and, rather, keep them proprietary. Hence, I was not able to build my work off of any existing products. In

a sense, this project aims to build something new.

Chapter 3

Methods

In this Section, I discuss my solution to the problem and various alternatives I considered along the way.

3.1   Overview

I solved this problem by formulating it as a search problem. By this I mean, given a schematic of a circuit, I start from an empty protoboard, and I search through the space of all possible protoboard layouts to find the protoboard corresponding to the schematic at hand. The space of all possible protoboards is very large (?), so I utilize various simplifications and heuristics to facilitate the search.

I broke down the problem into two parts. The first task is finding a placement of all the circuit pieces on the protoboard. The second task is wiring them up appropriately.

3.2   Part 1: Piece Placement

Let us first consider how to place a set of circuit pieces on the protoboard for a given circuit schematic. Any given circuit may contain resistors, Op Amps, pots, motors, head connector parts, or robot connector parts. For each of these components, we must put down a corresponding piece on the protoboard. As each piece may be placed on the protoboard in one of many different ways, I first decided on a fixed

Figure 3-1: Various acceptable ways of putting each of the circuit pieces on the protoboard.

set of allowed placements for each of the pieces. Figure 3-1 presents these acceptable placements. Resistors are placed in the middle strip of the protoboard. Op Amp pieces are also placed in the middle strip of the protoboard, but with two possible orientations. Op Amp pieces are unique in that each Op Amp pieces contains two Op Amps within it. Thus, we face the task of packaging the Op Amps in the schematic in the "best" possible way, i.e. so as to require as little work as possible when wiring the pieces together. Section 5.1 more precisely discusses the number of possibilities. Pots have two possible vertical positions (TODO: this is wrong!) as well as two possible orientations. The connector pieces have two possible vertical positions each.

3.2.1 Choosing a Placement

When choosing a placement of circuit pieces on the protoboard, we have at hand a plethora of options. First we must choose among a possibly large number of ways to package together the Op Amps in the circuit. For each possible packaging of Op Amps, we must consider various ways of placing the pieces on the protoboard, even with the restrictions put forth above.

18

Simplifications

I reduce this large number of options by only allowing placements in which no two pieces share a column. This is not necessary in general, but the number of pieces necessary for a typical 6.01 circuit would certainly fit in this framework.

Next, I specify that there be exactly two columns on the protoboard separating each consecuitive pair of pieces, unless the pieces are both resistors, in which case there must be exactly one column separating them. These numbers of columns were chosen to leave enough space for wiring. Given a set of pieces to be put on the protoboard, this specification reduces the problem of choosing a placement for the pieces to finding an *order* of the pieces together with choosing their respective vertical locations and orientations.

Given these simplifications, we have various options as to how to pick a placement.

Random Placement

One simple alternative may be to choose a placement randomly. That is, we choose an Op Amp packaging randomly; we choose an order of the pieces randomly; and we choose the vertical locations and orientations of the pieces randomly as well. The advantage of this approach is that it gives us a placement very quickly without requiring much computation. On the other hand, we may end up placing two pieces that need to be connected to each other very far apart, and we will have a difficult time doing the wiring. Hence, we ought to consider alternatives in which we take into account the task of wiring. We should try to place the pieces so as to require as little work during wiring as possible.

Minimal Heuristic Cost

The key idea is that if two pieces are meant to be connected together by wires, then they ought to be placed close to each other on the protoboard. We can capture this idea by assigning heuristic costs to the placements and choosing a placement that produces the minimal heuristic cost.

Let us first devise the cost function to achieve this goal. Given a circuit schematic and a corresponding placement of the circuit pieces on the protoboard, what do we need to connect with wires? Well, every pair of components in the schematic that are connected by a wire gives us a corresponding pair of locations on the protoboard that ought to be connected by wires. However, we can express this requirement a little bit more concisely. We ought to consider all of the nodes in the schematic, and find the circuit components in the schematic that are connected to the respective nodes. Now for each node in the circuit, we get a set of locations on the protoboard that ought to be interconnected. The first step in devising the cost function we are looking for is to have a way to estimate the cost of connecting two locations on the protoboard. A simple such cost function that comes to mind is the Manhattan distance between the two locations. Recall that we want to produce aesthetically pleasing protoboard layouts, and one of the requirements in achieving this goal is only using horizontal and vertical wires (i.e. no diagonal wires) so the Manhattan distance cost is appropriate. Given this heuristic cost for connecting two locations with wires, we can define the heuristic cost for interconnecting the locations associated with a particular node to be the weight of the minimum spanning tree of the locations. Now we can define the cost of a placement to be the sum over all nodes in the circuit of the cost for interconnecting the locations for each node.

Now that we have a cost function for placements, we can aim to find a placement with the minimal cost. However, this involves trying all possible orderings of the pieces with which we are working. For example, if we are trying to order 10 pieces, we would need to look at $10! = 3628800$ possible orderings (?). Note that this is in addition to searching over all possible ways of packaging the Op Amps together. It is clear to see that the search for a minimal cost placement quickly gets out of hand. So we aim to find a placement that has a very small, though maybe not minimal, cost.

Small Heuristic Cost

Algorithm 1 presents a polynomial-time procedure that orders a given list of pieces in a way that results in a small cost. The algorithm places one of the pieces at a time,

starting from an empty placement. It relies on two ideas. First, once a piece has been placed, all the pieces that are connected to it will be placed soon after so that it is more likely that those pieces are placed close to it. Second, we place the pieces with the most nodes first since those are the ones that most likely have connections with many other pieces.

Algorithm 1: Producing a circuit piece placement with small heuristic cost.

Data: A list P of circuit pieces.
Result: A list R of circuit pieces representing a placement.

Sort P in decreasing number of nodes on the respective pieces
$Q \leftarrow$ empty Queue
$R \leftarrow$ empty List
while *P is not empty* **do**
    Pop the first piece in P and push it onto Q
    while *Q is not empty* **do**
        $p \leftarrow Q$.pop()
        Consider all vertical locations and orientations of p
        Place p at an index in R that minimizes the cost of P
        foreach *piece q in P connected to p* **do**
            Pop q out of P and push it onto Q

Using one of the above methods, we can find a placement of circuit pieces on a protoboard. Our next task is wiring them together to produce a circuit equivalent to the circuit schematic of interest.

3.3   Part 2: Wiring

In the previous section we discussed what locations we need to wire together: for every node in the circuit, we get a set of locations on the protoboard that need to be interconnected. The question now is how to achieve this wiring. We approach the problem as a search problem and use the $A*$ search algorithm to solve it.

3.3.1   Using $A*$

When using the $A*$ algorithm, we need to design four things:

1. The notion of a vertex[1] in the search tree, the cost associated with a vertex, and how we obtain the neighbors of a vertex,

2. The starting vertex,

3. How we identify whether a particular vertex in the search tree achieves the goal of the search, and

4. A heuristic function that estimates the distance from a given vertex to a goal vertex.

3.3.2 Vertices

Each vertex will hold a representation of some protoboard. Each representation will contain all of the pieces, and possibly a set of wires interconnecting the pieces. The starting vertex will have all of the pieces but no wires.

We obtain the neighbors of a vertex by taking the current protoboard and producing new ones in which we place exactly one wire at various locations. We choose the starting point of a wire to be any one of the free locations on the protoboard that is already connected to one of the pieces, and we extend the wires in all possible vertical and horizontal directions up to some fixed wire length. Note that we need to take great care when placing wires in order not to short different nodes. We discard any vertices that arise from placing a wire that shorts two different nodes.

The way we define the cost of a vertex, i.e. the cost of getting from the starting vertex to a vertex of interest, depends on what we consider to be an aesthetically pleasing protoboard layout. In general, we want to penalize having long wires, many wires, or crossing wires. In my implementation, while I have a large penalty for two crossing wires of opposite orientations (i.e. vertical and horizontal), I do not allow crossing wires that have the same orientation as this configuration is particularly difficult to physically build and debug. Finally, we want to favor making a

---

[1]The prefered name is "node" but I will use vertex since we already use node to refer to nodes in circuits.

desired interconnection between locations on the protoboard. I chose my penalties experimentally (?).

Each vertex will not only hold a protoboard, but it will also hold a set of pairs of locations on the protoboard that need to be connected by wires. Each pair (loc_1, loc_2) of locations tells us that we need to have a set of wires connecting some location connected to loc_1 to some location connected to loc_2.

An important consideration we need to make is how we want to organize the search. Recall that we have a set of nodes in the circuit of interest, and for each node we have a set of locations that need to be interconnected. Given this information, we may choose one of the following three strategies to carryout the search:

1. For each node, collect a set of pairs of locations on the protoboard corresponding to a minimum spanning tree of the locations for that node, so that if all pairs of locations in this spanning tree are connected, then the locations for the node will be interconnected. Collect all such pairs of locations for all of the nodes in the circuit, and have the starting vertex hold this set of pairs of locations.

2. Treat each node separately. That is, iteratively interconnect the locations for each of the nodes until there are no more nodes in the circuit.

3. Treat each pair of locations that needs to be connected separately. That is, iteratively connect pairs of locations that need to be connected until there are no more pairs.

The choice of one of these strategies has a significant effect on the outcome of the search. We will discuss the difference in detail in Chapter 4.

3.3.3 Goal test

Given a vertex, we know that it is a goal vertex if all of the pairs of locations it holds are already connected in the protoboard it holds.

3.3.4   Search heuristic

In $A*$ search, choosing the right heuristic can often make the search much more efficient. In our problem, one option we have is not to use a heuristic, and that alternative will be explored in Chapter 4. However there is a natural heuristic that suggests itself that we ought to consider. Given a vertex, we can estimate its distance from a goal as follows. For each pair of locations (loc_1, loc_2) that need to be connected, we could consider its distance from the goal to be the smallest Manhattan distance between any location connected to loc_1 and any other location connected to loc_2. To compute the heuristic cost of a vertex, we simply add up this value for each of the pairs of locations that need to be connected. Chapter 4 presents the performance of this heuristic verses using no heuristic.

3.4   Combining the methods

With the methods discussed so far, we aimed to completly solve the layout problem with one placement method and one wiring method. However, as we will soon see, such an algorithm is bound to fail on some set of schematics. When we ultimately put the final algorithm in front of students, we would like to avoid failure. The algorithm should always be able to generate layouts. Generating a layout with a few diagonal or crossing wires is better than silently failing and leaving the student empty handed. Here, we will discuss how we combine the methods described so far into one layout algorithm. The motivation for this combination will be discussed in Chapter 5 based on the data we obtain for the alternatives described above. Algorithm 2 presents the combination algorithm.

There are a few key items in Algorithm 2. The algorithm works by attempting to solve the problem in four different ways: two different ways of doing placement together with two different orders of wiring pairs. Note, therefore, that we have chosen per-pair wiring. If any one of the four trials succeeds, the algorithm immediately returns the corresponding layout. If all four trials fail, on the other hand, the algorithm picks the "best" partial solution and completes that solution by placing possibly di-

---
Algorithm 2: Layout algorithm obtained by combining multiple ideas.

Data: A circuit schematic C.

Result: A layout corresponding to C.

foreach *Placement cost metric M in (DISTANCE, BLOCKING)* **do**
    $P \leftarrow$ Placement for C by using cost metric M.
    foreach *Order O in (INCREASING, DECREASING)* **do**
        pairs $\leftarrow$ Pairs of location on P to connect given scematic C and
        connection order O.
        foreach (loc_1, loc_2) *in pairs* **do**
            Attempt to connect loc_1 and loc_2 on P.
            If sucessful, update P accordingly and then post-process P.
            If not, record that the pair (loc_1, loc_2) was not successfully
            connected.
        If all pairs are successfully connected, post-process and return P.

Pick "best" unfinished layout.
Connect remaining pairs with shortest possible wires (possibly diagonal).
Post-process and return resulting layout.

---

agaonal wires, "best" being the partial solution that would require the least amount of additional wiring. This last step makes it highly unlikely that the algorithm will ever fail. The only way for the algorithm to fail is for there to be two nodes that need to be connected where all of the protoboard locations for at least one of the nodes is occupied, which is highly unlikely to happen. This high success rate comes at the cost of placing wires that will almost surely cross other wires on the board.

The algorithm also has a post-processing step that attempts to to improve the layout. The post-processing step makes three types of simple changes to the layout. First, we throw away any wires that do not serve to connect two parts of the circuit. That is, we throw away wires which have one end connected to the circuit, and the other end connected to nothing at all. This could happen in the search done by the wiring step, though very rarely. Second, we truncate long vertical wires into an equivalent set of smaller wires. For example, a wire going from one of the top rails to one of the bottom rails can be replaced by three smaller wires making the same connection. This change frees up rows for subsequent connections. Finally, if shifting a horizontal wire up or down results in fewer wire crosses, we make that change.

The last important aspect of this final algorithm not explicitly stated in Algorithm

2 is that, as it is the algorithm that will be put in front of students, it will be allowed to use wires of a select few lengths. The kits that students work with certainly do not come with wires of all lengths, so we force the wiring step to use wires of only those allowed lengths.

3.5   Treating Resistors as Wires

The discussion in Section 3.2 presented that we treat resistors just as we do the other components. That is, we give each resistor a fixed place on the protoboard in the first step of the algorithm before the wiring step. However, resistors have the special property among the circuit pieces that they can be thought of as wires of length 3. Hence, it may be possible to handle resistors in the wiring step instead of the placement step. Chapter 4 presents a comparison of these two approaches.

Note that treating resistors in the wiring step is not a trivial task. First, there may be nodes in the circuit that are connected to some resistors, but no other pieces. In this case, we must be sure to reserve space on the protoboard for that node as the wiring step relies on the presence of each node on the protoboard. Second, we must keep careful track of pairs of locations that need to be connected by simple wires and pairs of locations that need to be connected using resistors.

3.6   Evaluation Method

Here we present how we go about evaluating a particular solution to the problem. How can we tell if a layout tool is good? In particular, how can we tell if a layout tool is good enough for the purposes of 6.01 labs. To answer these questions well, we need to test the layout tool on numerous schematics and analyze its performance on laying out those schematics. As manually generating numerous test schematics is tedious and very time-consuming, we devised a method to randomly generate thousands of test schematics.

The random schematic generation goes as follows. We create 6 basic sub-parts

of schematics. These 6 bases are depicted in Figure 3-2. These bases cover all of the components that may be necessary in a 6.01 circuit. Each sub-part also offers at least 3 points of connection with other subparts. The random generation algorithm takes all possible combinations of 6 bases, allowing for repetition of bases with some restictions. The Head Connector and Robot Connector bases can appear at most once as there is no need for more than one of each of these in 6.01 labs. The pot-follower base (i.e. the base that contains one pot and a follower op amp) can appear at most twice as we never need more than two pots in 6.01 circuits. The motor base can also appear at most twice as we never need more than two motors per circuit in 6.01 labs. The other bases, T-resistor configuration and voltage divider, can be repeated up to 6 times. For a given combination of bases, we generate up n schematics, where n dependes on the number of bases in the configurations (larger for configurations with more bases in them). We generate the n schematics by choosing a number of interconnections between the bases from 0 to $n - 1$. For each number of interconnections, we generate a schematic in which we put that many randomly chosen wires interconnecting the bases in the combination. Figure 3-3 presents a sample randomly generated schematic.

This scheme produces a total of 4425 test schematics out of a possible total of approximately $1.2e27$. When testing a particular algorithm on these test schematics, we run the algorithm on each test schematic 10 times. Chapter 4 presents the data collected in this manner comparing the various alternatives discussed in this Chapter.
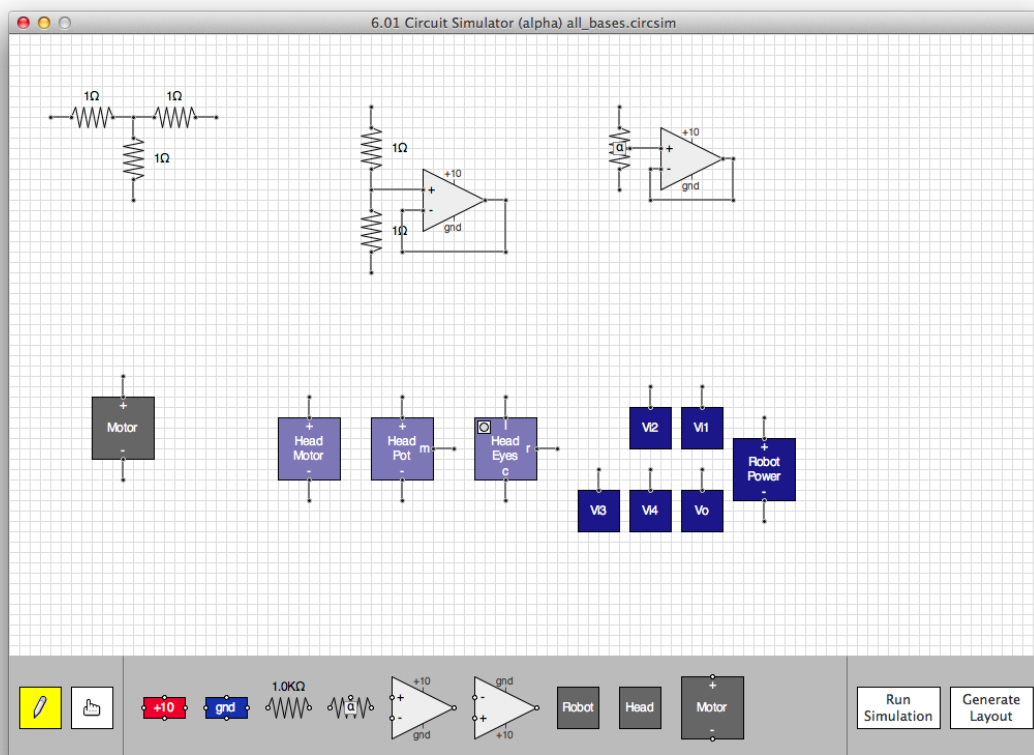
Figure 3-2: Bases for random schematic generation.

Figure 3-3: Sample randomly generated schematic.

Chapter 4

Results

In Chapter 3 we discussed a general solution to the protoboard layout problem, and various alternatives that can be used in implementing the solution. Figure 4 presents the alternatives in a structured way. Here, we will explore these alternatives and compare them quantitatively. This section will provide data that is useful in comparing the alternatives, and the data will be discussed in detail in Chapter 5.

Figure 4-1: All possible alternatives to the algorithm.

As comparing all 40 possible implementations of the algorithm is tedious, we will compare the alternatives for each aspect of the algorithm while holding other aspects fixed. Hence, we will carryout the following comparisons:

1. Placement: Distance vs. Blocking. Wiring method will be per-pair, decreasing,

resistors will be treated as components, and we will use $A*$.

2. Wiring: All pairs vs. Per-node, decreasing vs. Per-node, increasing vs. Per-pair, decreasing vs. Per-pair, increasing. Placement method will be blocking, resistors will be treated as components, and we will use $A*$.

3. Resistor treatment: As components vs. As wires. Placement method will be blocking, wiring method will be per-pair, decreasing, and we will use $A*$.

4. Search: $A*$ vs. Best First. Placement method will be blocking, wiring method will be per-pair, decreasing, and resistors will be treated as components.

The data to compare the alternatives is gathered as described in Chapter 3. We run the algorithm on 4425 randomly generated schematics of varying complexities. The algorithm is run 10 times on each schematic.

In comparing alternatives, there are 3 items we will consider:

1. Which alternative is the most successful?

2. Which alternative, when successful, takes the least amount of time?

3. Which alternative, when successful, produces the best layouts?

In comparing success, we will look at bar graphs of the number of successes on each of the 4425 schematic out of the 10 runs. We will also look at tables that provide the same data in more detail. To get an understanding of how the success rates vary with complexity, we will look at plots of circuit complexity versus success rate, where our measure of circuit complexity will be the number of pins in the circuit. A pin in a circuit is a connection point on a circuit component that is connected by wires to another connection point (on the same component or a different component). Hence, the number of pins in the circuit is a reasonable metric for how complex the circuit is. Figure 4-2 presents a histogram of the number of pins in the schematics in dataset that will be used to do all comparisons in this Chapter.

31

Figure 4-2: TODO.

In comparing success time, we will look at CPU time spent on the wiring step, as the placement step has much less variability. We will look at plots of circuit complexity versus wiring times to do the comparisons.

In comparing the goodness of layouts, we will compare numbers of wires, total lengths of wires, and numbers of wire crosses as functions of circuit complexity.

Below we will present the data for each of the comparisons outlined above. Note that in all figures that follow, error bars indicate 1.96 times the standard error.

Discussion of these results will follow in Chapter 5.

4.1 Comparing placement methods
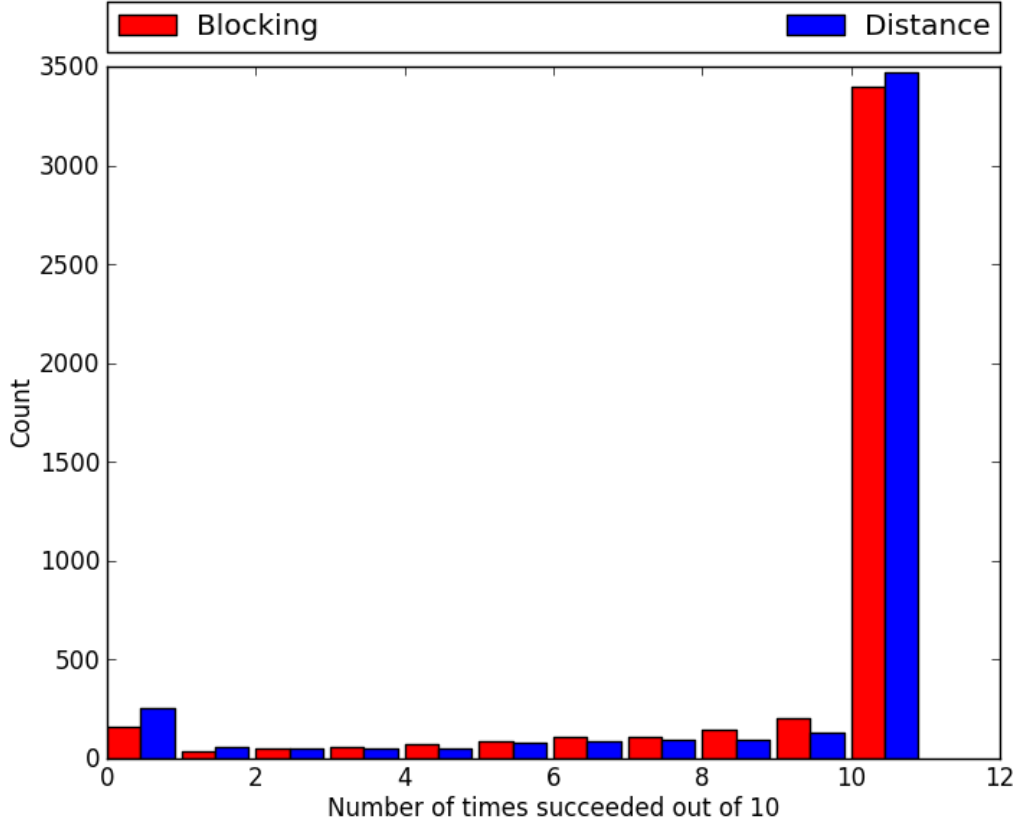
Figure 4-3: Placement method comparison: success rates.

	Number of times succeeded out of 10										
	0	1	2	3	4	5	6	7	8	9	10
Blocking	162	38	51	57	72	85	109	106	144	203	3398
	0.04	0.01	0.01	0.01	0.02	0.02	0.02	0.02	0.03	0.05	0.77
Distance	258	55	54	50	52	77	86	97	93	130	3473
	0.06	0.01	0.01	0.01	0.01	0.02	0.02	0.02	0.02	0.03	0.78

Table 4.1: Placement method comparison: success rates.

Figure 4-4: Placement method comparison: success rate trends.

Figure 4-5: Placement method comparison: wiring time trends.

Figure 4-6: Placement method comparison: layout quality trends.

4.2 Comparing wiring methods

Figure 4-7: Wiring method comparison: success rates.

	Number of times succeeded out of 10										
	0	1	2	3	4	5	6	7	8	9	10
All	458	114	111	112	127	145	177	139	172	227	2643
	0.10	0.03	0.03	0.03	0.03	0.03	0.04	0.03	0.04	0.05	0.60
Node D	195	50	58	66	104	83	125	176	162	268	3138
	0.04	0.01	0.01	0.01	0.02	0.02	0.03	0.04	0.04	0.06	0.71
Node I	154	50	55	62	85	71	106	141	156	217	3328
	0.03	0.01	0.01	0.01	0.02	0.02	0.02	0.03	0.04	0.05	0.75
Pair D	162	38	51	57	72	85	109	106	144	203	3398
	0.04	0.01	0.01	0.01	0.02	0.02	0.02	0.02	0.03	0.05	0.77
Pair I	177	40	59	54	91	92	100	118	132	212	3350
	0.04	0.01	0.01	0.01	0.02	0.02	0.02	0.03	0.03	0.05	0.76

Table 4.2: Wiring method comparison: success rates.

Figure 4-8: Wiring method comparison: success rate trends.

Figure 4-9: Wiring method comparison: wiring time trends.

Figure 4-10: Wiring method comparison: layout quality trends.
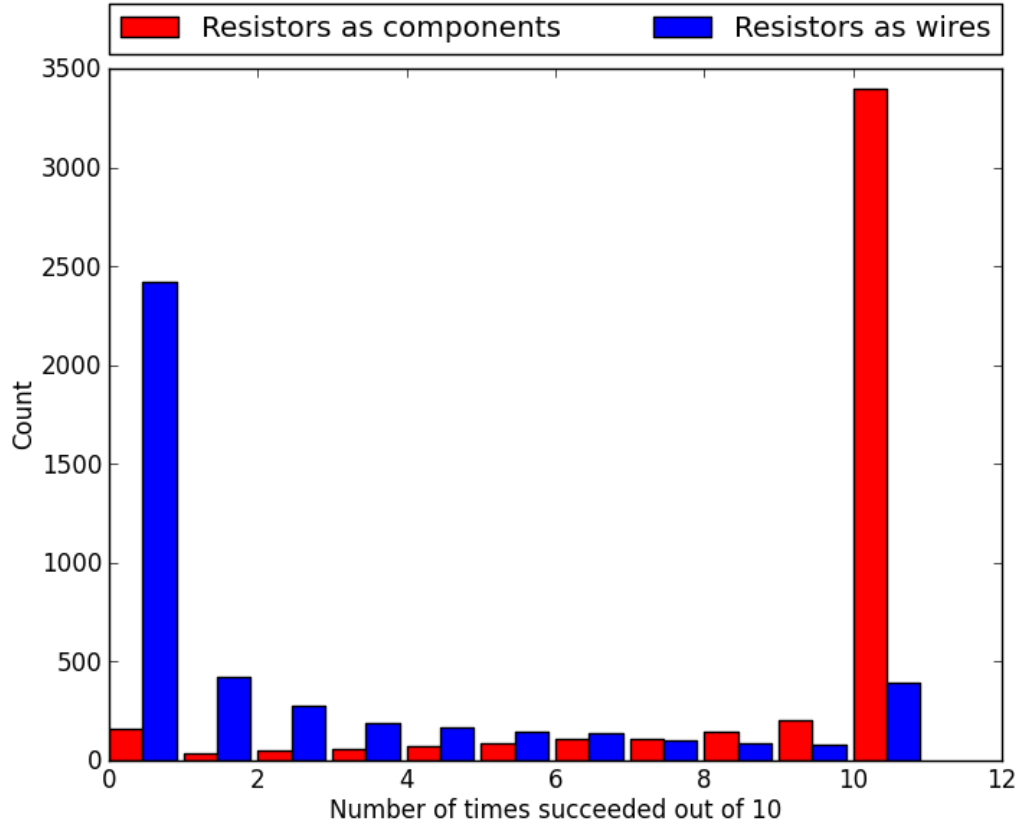
4.3 Comparing resistor treatments

Figure 4-11: Resistor treatment comparison: success rates.

	Number of times succeeded out of 10										
	0	1	2	3	4	5	6	7	8	9	10
As components	162	38	51	57	72	85	109	106	144	203	3398
	0.04	0.01	0.01	0.01	0.02	0.02	0.02	0.02	0.03	0.05	0.77
As wires	2424	421	276	187	171	146	140	101	91	77	391
	0.55	0.10	0.06	0.04	0.04	0.03	0.03	0.02	0.02	0.02	0.09

Table 4.3: Resistor treatment comparison: success rates.

Figure 4-12: Resistor treatment comparison: success rate trends.

Figure 4-13: Resistor treatment comparison: wiring time trends.

Figure 4-14: Resistor treatment comparison: layout quality trends.

4.4 Comparing search methods

Figure 4-15: Search method comparison: success rates.

	Number of times succeeded out of 10										
	0	1	2	3	4	5	6	7	8	9	10
$A*$	162	38	51	57	72	85	109	106	144	203	3398
	0.04	0.01	0.01	0.01	0.02	0.02	0.02	0.02	0.03	0.05	0.77
Best First	6	5	2	1	13	10	29	45	84	245	3985
	0.00	0.00	0.00	0.00	0.00	0.00	0.01	0.01	0.02	0.06	0.90

Table 4.4: Search method comparison: success rates.

Figure 4-16: Search method comparison: success rate trends.

Figure 4-17: Search method comparison: wiring time trends.

4.5 Combined algorithm

Here we provide data for the combined algorithm presented in Section 3.4. To generate this data, we use a different dataset of 4425 schematics. As desired, the algorithm has a 100% success rate. Figure 4-19 gives a breakdown of how the algorithm succeeded. The first four columns correspond to success from one of the four combinations of placement and wiring methods. The last 5 columns correspond to layouts in which none of the four combinations was successful and we had to force a certain number of wires. Figure 4-20 gives the average number of forced wires as a function of circuit complexity. Figure 4-21 gives the average total time taken by the algorithm as a function of circuit complexity. Finally, Figure 4-22 gives statistics on the quality of

Figure 4-18: Search method comparison: layout quality trends.

the layouts produced as a function of circuit complexity.


Figure 4-19: TODO.

Figure 4-20: TODO.

Figure 4-21: TODO.

Figure 4-22: TODO.

Chapter 5

Discussion

In this section we provide justifications for the choices made in solving the protoboard layout problem, and also detailed analysis of the data presented in Chapter 4.

5.1 Justifying Placement Choices

Resistors

For the sake of simplicity, and to significantly reduce the search space (?), for every resistor in the schematic, I use one resistor piece on the protoboard placed in the middle strip of the protoboard as shown in Figure 3-1. This choice, i.e. allowing the resistor pieces to only reside in the middle strip of the protoboard, is critical as the resistor pieces can generally be placed at numerous places on the protoboard. With this restriction, there are 63 slots available for one resistor. Without this restriction, there are a total of 763 slots available. The restirction is good when we consider the reduction in the search space size. On the other hand, the restriction is bad when we consider the size of circuits the algorithm can layout. Given that the number of resistors in a typical 6.01 circuit is very small (?), this restriction proves to be very useful.

n	$f(n)$
1	1
2	3
3	7
4	25
5	81
6	331
7	1303
8	5937
9	26785
10	133651

Table 5.1: Number of ways of packaging together n Op Amps for various values of n.

Op Amps

Op Amps are the trickiest components to handle because each Op Amp package put on the protoboard contains two Op Amps within it. Equation 5.1 presents an expression for the value $f(n)$, the number of different ways to package together n Op Amps. For example, if we have 2 Op Amps, we can either use one Op Amp package for each, or put them both in the same package, which we can do in one of two different ways. Hence, $f(2) = 3$. To get a sense of how many different packagings are possible, Table 5.1 gives the values of $f(n)$ for various n.

$$f(n) = \sum_{k=0}^{\lfloor \frac{n}{2} \rfloor} \frac{n!}{k!(n-2k)!} \tag{5.1}$$

Our placement approach explores all possible ways of packaging up the Op Amps. We do this because the typical 6.01 circuit contains no more than 6 Op Amps, and so we are tasked with exploring at most 331 alternatives, which doesn't happen to be too computationally intensive. If the algorithm was meant to handle a larger number of Op Amps (for example 10), then this approach would be silly, and we would have to resort to perhaps considering the packagings that require the fewest number of Op Amps to make a choice.

Pots

Each pot piece can be placed in one of two vertical locations on the protoboard. Each pot piece also has two possible orientations. TODO: really 6 different vertical locations.

Head, Motor, and Robot Connectors

Our choices for the connector pieces is certainly one of the two presented in Figure 3-1. Any other placement of the connector pieces would unecessarily take up valueable rows.

5.2 Explaining the Results

Chapter 4 presented quantitative data to compare the various alternatives we have in solving the protoboard layout problem. Here, we will analyze that data and give reasonings for why we obtained the results that we obtained. For each comparison, we present exemplar layouts generated by the alternative methods for the schematic shown in Figure 5-1.

5.2.1 Comparing placement methods

In Figure 4-3 and Table 4.1 we see that the distance method exceeds the blocking method in number of circuits solved 10 times out of 10 as well as in the number of circuits solved 0 to 2 times out of 10. On the other hand, the blocking method exceeds the distance method in number of circuits solved from 3 to 9 times out of 10. Overall, the blocking placement method is slightly more successful than the distance placement method. Despite this difference, we note that the two alternatives have very close success rates. As a function of circuit complexity, Figure 4-4 suggests that the two alternatives have almost identical success rates. As we would expect, we certainly observe that as the complexity of the circuits increases, success rate generally decreases for both of the placement methods. The curves in Figure 4-4
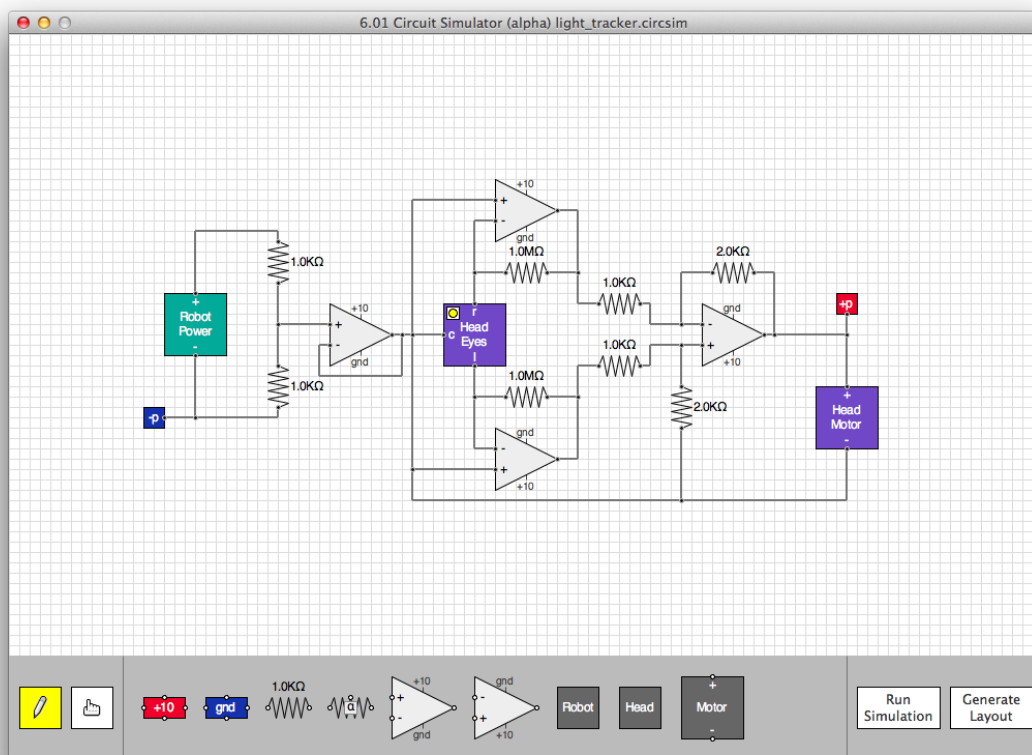
Figure 5-1: Exemplar schematic.

seem to shoot back up at the far end of the figure, but this is a result of the fact that there are very few circuits at that end of the figure, on which the algorithm happened to be consistently successful.

When we consider wiring time as the basis for comparison, we observe from Figure 4-5 that, once again, the two metods are very similar. We see that the distance method generally takes slightly less wiring time than the blocking method, but the difference between the two is almost negligible. As we would expect, we certainly see that as the complexity of the circuits increases, the amount of time spent by the wiring step also increases. As we did for the success rate trends as a function of circuit complexity, we observe that there are outliers at the far end of the figure due to a very small sample of the most complex circuits in the randomly generated schematic dataset.

Finally, let us look at layout complexity as the basis for comparison. Figure 4-6 presents graphs that compare numbers of wires, numbers of wire crosses, and total wire lengths. We first observe that the number of wires used by the two methods are almost identical. As the complexity of the circuit increases, we see that the blocking method uses more wires than does the distance method, but by and large the values are very comparable. This makes intuitive sense as we are rarely required to use more wires than absolutely necessary (keeping wires horizontal and vertical) to connect a pair of locations (?). When we look at the number of wire crosses in the layouts, we see that the blocking method consistently results in more wire crosses. Similarly, when we look at the total length of wires used in the layouts, the blocking method exceeds the distance method consistently, with the difference getting higher as complexity increases. This can be explained by the fact that the blocking method may require elaborate connections, especially as the circuits get more complex while the distance method is tuned to require wirings between pairs of locations that are closer together.

It is difficult to conclusively pick the better placement method from these results. It is clear that both methods have their strengths and weaknesses. In Section 5.2.5 we will discuss how combining these two methods can get us the best of both worlds (?).

Figure 5-2: Blocking vs. Distance.

5.2.2   Comparing wiring methods

To compare the wiring methods, let us start with the success rate for each method. Figures 4-7 and 4-8, and Table 4-7 provide the appropriate data. The very first fact we observe is that the all-pairs method has a much smaller success rate than all of the other alternatives, especially as circuit complexity increases. The reason for this, in large part, is the fact that we impose a cutoff on the number of nodes to expand in the search (TODO: this must be discussed in Methods). As we are using just the one search to connect all pairs of locations, the search cutoff would be expected to have more of an effect on the all-pairs method. The other four alternatives have very comparable success rates.

Next we look at Figure 4-9 to compare wiring times for the five methods. Once again, we observe that the all-pairs methods takes significanlty more time than the other methods. When attempting to connect all pairs in one search, the method searches for an appealing layout, which may require searching through a large number of alternatives, especially for more complex circuits. Hence, we would expect the all-pairs methods to generally take more time than the other alternatives. We also observe from Figure 4-9 that the wiring times for the two per-node methods are comparable, and that the wiring times for the two per-pair methods are also comparable, but that the per-node wiring times are generally bigger than the per-pair wiring times. This trend is also expected as the per-node methods attempt to connect multiple pairs of locations at once while keeping the layout pleasing, and this generally requires searching through more alternatives than connecting each of the pairs of locations individually. It is important to note that per-node, increaseing generally takes more time than per-pair, decreasing, and also that per-pair, increasing generally takes more time than per-pair, decreasing.

Finally let us look at Figure 4-10 to compare the quality of the layouts produced by the five alternative wiring methods. First, as we observed in the placement method comparison, we see that there is very little difference in terms of number of wires used and the total wire length. However, there are significant differences in the number

of wire crosses. We see that the all-pairs method generates layouts with much fewer wire crosses than the other methods. This is completely expected since we run one search to connect all pairs of locations while attempting to keep the layout as nice as possible. Conversely, the per-pair, decreasing and per-node, decreasing methods result in the most number of wire crosses. It is very interesting to note that the per-node, decreasing method produces more wire crosses on average than the per-pair, increasing method. Here we observe that the order in which we consider pairs of locations has a telling effect on how good the layouts will be. In essence, connecting the harder pairs of locations generally produces more wire corsses.

What we have observed is that while the all-pairs method is the least successful method and the one that generally takes the longest among the five, it generally produces the best layouts when it does succeed. On the other hand, the alternatives that break down the problem into smaller pieces succeed more often and finish more quickly while producing worse results. Furthermore, the more finely we breakdown the problem, the faster the overall algorithm. Lastly, ordering subproblems from hardest to easiest has the effect of making the overall wiring step complete faster but producing worse results than the reverse order and not necessarily getting a markedly better success rate.

5.2.3   Comparing resistor treatments

We now look at the two alternative treatments for wires: as components vs. as wires. From Figures 4-11 and 4-12 and Table 4.3, we observe that treating resistors as wires results in a considerably smaller success rate than treating resistors as components. Recall that the wiring method used here is per-pair, decreasing wiring. When treating resistors as wires, we get more location pairs to connect. The key difference for the location pairs that correspond to resistors is that in connecting the given pair of locations with wires, one of the wires is required to have length 3, the length of a resistor. Intuitively, this restriction should make the search harder, and the data we have obtained suggests that the search is indeed significantly more difficult with the restriction than without.

60

Figure 5-3: All-pairs vs. Per-node, increasing vs. Per-node, decreasing vs. Per-pair, increasing vs. Per-pair, decreasing.

When we treat resistors as wires, and the search does succeed, we observe from Figure 4-13 that it takes more time than treating resistors as wires, as expected.

Finally we look at the quality of layouts produced by the two methods. First, we observe from Figure 4-14 that treating resistors as wires produces layouts with fewer wires and smaller total length of wires than treating resistors as components. This is certainly expected because each resistor can be suitably placed to require fewer wires than would be required if we forced the resistor to reside in the middle strip of the protoboard. We do note, however, that treating resistors as wires results in more wire crosses despite requiring fewer wires, especially as circuit complexity increases. This can be attributed to the fact that treating resistors as wires is likely to produce more compact layouts in which wire crosses are more likely to happen than the spreadout settings we force when we treat resistors as components.

In this comparison, we can easily state that, with a limited search (i.e. with a search in which we have a limited number of states to expand), treating resistors as components is a better strategy than treating resistors as wires based on the big difference in success rate.

5.2.4 Comparing search methods

We now compare our two alternative search algorithms: $A*$ and Best First Search. Let us start by comparing success rates. Figures 4-15 and 4-16 and Table 4-15 present the appropriate data. We clearly observe that Best First Search is much more successful than $A*$. 98% of the test circuits were solved at least 8 times out of 10 when we used Best First Search, versus 85% when we used $A*$. This result is not surprising because Best First Search hungrily hunts for layouts that satisfy the connection requirements without caring for the aesthetics of the layouts. Hence, Best First Search is much less succeptible to the 300 states to expand restriction than $A*$. The fact that the quality of the results we get from Best First Search are worse is clearly evident from Figure 4-18. Most importantly, the number of wire crosses in the layouts produced by Best First Search are markedly greater than the number of wire crosses in the layouts produced by $A*$. We also observe that the total wire length is greater when

Figure 5-4: Resistors as components vs. Resistors as wires.

Figure 5-5: $A*$ vs. Best First Search.

using Best First Search. The fact that Best First Search settles for any layout that satisfies the connection requirements suggests that it should finish more quickly in addition to being more successful. Firgure 4-17 supports exactly this expectation.

Our choice of a search algorithm forces us to consider a treadoff between quick success and quality. If we choose Best First Search, most runs will be successful and terminate quickly, but will produce very poor results. If we choose $A*$, not as many runs will be successful, and the successful runs will take longer to terminate, but we will get much better layouts.

5.2.5   Putting them all together

Now we will discuss whey the combination algorithm given in Algorithm 2 is structured the way it is, and we will also discuss the data we obtained for the combined algorithm. Recall that the combined algorithm makes four attempts at generating a complete layout. We try both placement methods, the distance method first and the blocking method second. We try the distance method first because it tends to generate layouts with fewer crossing wires and smaller total wire length. We use per-pair wiring, and consider both orders of doing the wiring, increasing order first and decreasing order second. We use per-pair wiring because it takes considerably less time than either per-node or all-pair wiring, and neither of the other wiring methods has a better success rate. We try increasing order first because that tends to generate layouts with fewer crossing wires then decreasing order.

Let us now consider the data we obtained for the combined algorithm. Firstly, we saw that the combined algorithm had a 100% success rate, which is critical when we consider the fact that students will be using a tool that almost never fails. Next, Figure 4-19 presents very encouraging results as well. Not only is the algorithm 100% successful, on the test dataset of 4425 schematics, with 10 runs carried out on each schematic, the algorithm succeeded 98.2% of the time without having to put down forced wires. The algorithm was required to put down more than 2 forced wires only 0.1% of the time. When we look at Figure 4-20, we observe, as expected, that the number of forced wires increases with circuit complexity. From Figure 4-21 we see that as the circuit gets more complex, the amount of time the algorithm takes sharply increases. This is due to a more difficult placement task (especially as we have more Op Amps) as well as a more difficult wiring task. An important aspect of this Figure is that the maximum point on the plot occurs at less than 200 seconds. This indicates that on the test dataset, the algorithm took at most just over 3 minutes to run. This is encouraging from a practical stand point because 3 minutes is not terribly long for a student to have to wait for a layout to be generated. Finally Figure 4-22 presents trends of quality as a function of circuit complexity. Most importatnly, we observe

that the average number of wire crosses never exceeds 1.

5.3   Remarks

Why are these results encouraging? What are their implications? Relate back to Introduction to Thesis. What could have been done differently?

Appendix A

Schematic Drawing GUI

Discuss the features and capabilities of the GUI.