

Automatic Protoboard Layout

by

Michael Mekonnen

B.S. EECS, Massachusetts Institute of Technology (2013)

B.S. Mathematics, Massachusetts Institute of Technology (2013)

Submitted to the Department of Electrical Engineering and Computer
Science

in partial fulfillment of the requirements for the degree of

Masters of Engineering in Electrical Engineering and Computer Science

at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

February 2014

© Massachusetts Institute of Technology 2014. All rights reserved.

Author
Department of Electrical Engineering and Computer Science
November 25, 2013

Certified by
Dennis M. Freeman
Professor
Thesis Supervisor

Certified by
Adam J. Hartz
Lecturer
Thesis Supervisor

Accepted by
Professor Albert R. Meyer
Chairman, Department Committee on Graduate Theses

Automatic Protoboard Layout

by

Michael Mekonnen

Submitted to the Department of Electrical Engineering and Computer Science
on November 25, 2013, in partial fulfillment of the
requirements for the degree of
Masters of Engineering in Electrical Engineering and Computer Science

Abstract

As an important component of the Circuits module of the first Introduction to Electrical Engineering and Computer Science course at MIT (6.01), students design and build several circuits over the course of three weeks. When working on the more intricate circuits, an unfortunately large proportion of students' lab time is spent on laying out the circuits on protoboards. This project introduces a new circuit schematic entry tool for 6.01 capable of automatically generating protoboard layouts for circuits that students may design in the course. The tool allows for students to easily build and analyze circuit schematics through an intuitive graphical user interface, and automatically generates protoboard layouts that are almost always easy to build and debug. The layout problem is solved by utilizing the *A** search algorithm exactly as presented in 6.01.

Thesis Supervisor: Dennis M. Freeman
Title: *Professor*

Thesis Supervisor: Adam J. Hartz
Title: *Lecturer*

Acknowledgments

First and foremost, I would like to thank Professor Freeman and Mr. Hartz for their nonstop support and encouragement from the very start of this project to its finish. They were always happy to talk with me, listen to my ideas, and give me constructive feedback whenever I asked for it. I am very happy to have worked with and developed a lasting relationship with great mentors.

Next, I would be remiss not to express my gratitude to my always loving and supportive family. I would not have been able to complete this project without the constant encouragement from my parents and siblings.

Finally, I would like to thank the Fall, 2013 members of the 6.01 staff, including Instructors, Teaching Assistants, and Lab Assistants, for testing the tool and kindly giving me ideas on how to improve it.

Contents

1	Introduction	10
1.1	Problem Statement	10
1.2	Outline	11
2	Background	12
2.1	Technical Background	12
2.1.1	Circuit Components	12
2.1.2	Circuit Schematic	13
2.1.3	Protoboard	14
2.1.4	Protoboard Layout	14
2.2	Previous Work	16
2.2.1	CMax	16
2.2.2	Current work in automatic layout	18
3	Methods	20
3.1	GUI	20
3.2	Solving the Layout Problem	20
3.2.1	Part 1: Placement	22
3.2.2	Part 2: Wiring	28
3.2.3	Combining the Methods	32
3.2.4	Evaluation	34
4	Results	38

4.1	Comparing Placement Methods	42
4.2	Comparing Wiring Methods	48
4.3	Effect of limiting the number of vertices to expand in A^*	56
4.4	Comparing Search Methods	62
4.5	Combined Algorithm	67
5	Discussion	73
5.1	Search Space Size	73
5.2	Justifying Placement Choices	74
5.3	Explaining the Results	75
5.3.1	Comparing Placement Methods	75
5.3.2	Comparing Wiring Methods	77
5.3.3	Comparing Search Methods	79
5.3.4	Putting Them All Together	80
5.4	Further Work	81
5.4.1	Treating Resistors as Wires	81
5.4.2	Building Layouts Similar to Previously Generated Layouts . .	82
5.5	Remarks	82
A	Schematic Drawing GUI	83
A.1	Palette	83
A.2	Board	85
A.3	Analysis	86
A.3.1	Simulation	86
A.3.2	Layout	86
A.4	Other Features	88
A.5	Shortcuts	88

List of Figures

2-1	6.01 robot.	13
2-2	Sample circuit schematic: head angular position controller.	14
2-3	Physical protoboard. In the top and bottom groups, the dots are interconnected horizontally. In the middle two groups, the dots are interconnected vertically.	15
2-4	Protoboard layout for the schematic in Figure 2-2.	15
2-5	CMax layout for the schematic in Figure 2-2.	17
3-1	Sample schematic drawn on the schematic entry tool.	21
3-2	Various acceptable ways of putting each of the circuit pieces on the protoboard.	22
3-3	Two possible placements for the schematic shown in Figure 3-1. . . .	24
3-4	Placements to demonstrate cost functions.	26
3-5	Bases for random schematic generation.	35
3-6	Sample randomly generated schematic.	36
4-1	All possible alternatives to the algorithm.	38
4-2	Histogram of the complexity the 4425 schematics used for evaluation.	40
4-3	Exemplar schematic.	41
4-4	Blocking exemplar.	42
4-5	Distance exemplar.	42
4-6	Placement method comparison: success rates.	43
4-7	Placement method comparison: success rate trends.	44
4-8	Placement method comparison: wiring time trends.	45

4-9 Placement method comparison: layout quality trends.	46
4-10 Placement method comparison: layout badness trends.	47
4-11 All pairs exemplar.	48
4-12 Per-node (increasing) exemplar.	48
4-13 Per-node (decreasing) exemplar. We used the distance based placement method instead of the blocking based placement method to generate this exemplar as the combination of the blocking based method with this wiring method consistently failed on the schematic shown in Figure 4-3.	49
4-14 Per-pair (increasing) exemplar.	49
4-15 Per-pair (decreasing) exemplar.	50
4-16 Straight wiring exemplar.	50
4-17 Wiring method comparison: success rates.	51
4-18 Wiring method comparison: success rate trends.	53
4-19 Wiring method comparison: wiring time trends.	54
4-20 Wiring method comparison: layout quality trends.	55
4-21 Wiring method comparison: layout badness trends.	56
4-22 Histogram of number of vertices expanded for all pairs wiring.	57
4-23 Histogram of maximum number of vertices expanded for per-node (in- creasing) wiring.	58
4-24 Histogram of maximum number of vertices expanded for per-node (de- creasing) wiring.	59
4-25 Histogram of maximum number of vertices expanded for per-pair (in- creasing) wiring.	60
4-26 Histogram of maximum number of vertices expanded for per-pair (de- creasing) wiring.	61
4-27 A* Search exemplar.	62
4-28 Best First Search exemplar.	62
4-29 Search method comparison: success rates.	63
4-30 Search method comparison: success rate trends.	64

4-31	Search method comparison: wiring time trends.	65
4-32	Search method comparison: layout quality trends.	66
4-33	Search method comparison: layout badness trends.	67
4-34	Combined algorithm exemplar.	68
4-35	Combined algorithm success summary.	69
4-36	Combined algorithm total CPU time trend as a function of circuit complexity.	70
4-37	Combined algorithm quality trend.	71
4-38	Combined algorithm layout badness trend.	72
A-1	Parts of the schematic entry GUI.	84
A-2	Grouped components: Robot Connector and Head Connector.	84
A-3	Wire highlighting example.	87
A-4	Component highlighting example.	87

List of Tables

4.1	Placement method comparison: success rates.	43
4.2	Wiring method comparison: success rates.	52
4.3	Search method comparison: success rates.	63
5.1	Number of ways of packaging together n op-amps for various values of n	75
A.1	Shortcuts.	89

Chapter 1

Introduction

1.1 Problem Statement

In this paper we discuss the problem of automatic protoboard layout generation. Importantly, we are interested in automatically generating layouts that are easy to build, easy to debug, and aesthetically pleasing. The tool this paper discusses is strictly geared towards circuits that students would build in the Introduction to Electrical Engineering and Computer Science I[1] course at MIT (also known as 6.01).

We are motivated to solve the layout problem in order to let 6.01 students spend more of their lab time thinking about how to design circuits and less time thinking about how to lay them out on protoboards. In this project we built¹ a tool that lets students easily build and simulate circuit schematics through an intuitive graphical user interface (GUI). After building and testing a circuit schematic with the tool, a student can proceed to building the circuit on a physical protoboard based on the layout generated by the tool. Not only does the tool generate a layout, but it also makes it easy to relate the original schematic drawn by the student to the generated layout. With this tool, a student’s lab time will be spent mostly on designing, building, and testing circuits rather than on the difficult and arguably less instructive task of layout.

¹The tool was implemented entirely in the Python programming language.

1.2 Outline

Chapter 2 describes in detail the terminology that will be used in this paper and explores the current infrastructure available for 6.01 students as well as previous work done in automatic layout generation. Chapter 3 discusses how we solved the problem, including various alternatives considered in each part of the solution, and how we evaluated our solution. Chapter 4 presents data to compare the various alternatives discussed in Chapter 3 and also evaluates the final algorithm on a large test dataset. Finally, Chapter 5 presents arguments for the choices made in our solution to the problem and also elaborates upon the results presented in Chapter 4.

Chapter 2

Background

In this chapter we discuss essential background information to this project. First, we discuss the specific terminology used in this paper. Next, we discuss previous work done that relates to this project.

2.1 Technical Background

As we are producing a new teaching tool for 6.01 in this project, let us first discuss the scope of circuits in 6.01.

2.1.1 Circuit Components

The rudimentary circuit components used in 6.01 are resistors, operational amplifiers (op-amps), and potentiometers (pots). In addition to these basic parts, students may build circuits to control LEGO motors or to control aspects of robots designed specifically for 6.01. One of the 6.01 robots is depicted in Figure 2-1. The robots can be equipped with heads that contain three parts held together by a shaft: a LEGO motor, a potentiometer, and a plate containing two photosensors. The robot in Figure 2-1 has a head attached. To connect a layout to a LEGO motor, a student would use a 6-pin connector, and to connect a layout to a robot or a robot head, a student use an 8-pin connector.



Figure 2-1: 6.01 robot.

All together, the circuit pieces that may be used in constructing protoboard layouts in 6.01 are resistors, op-amps, pots, (6-pin) motor connectors, (8-pin) robot connectors, and (8-pin) head connectors.

2.1.2 Circuit Schematic

Throughout this paper, the term *circuit schematic* will refer to a drawing or a sketch of a circuit containing its components and all the interconnections between the components drawn as wires. This is what one would sketch on a piece of paper in the process of designing a circuit. Figure 2-2 presents an example of a circuit schematic.

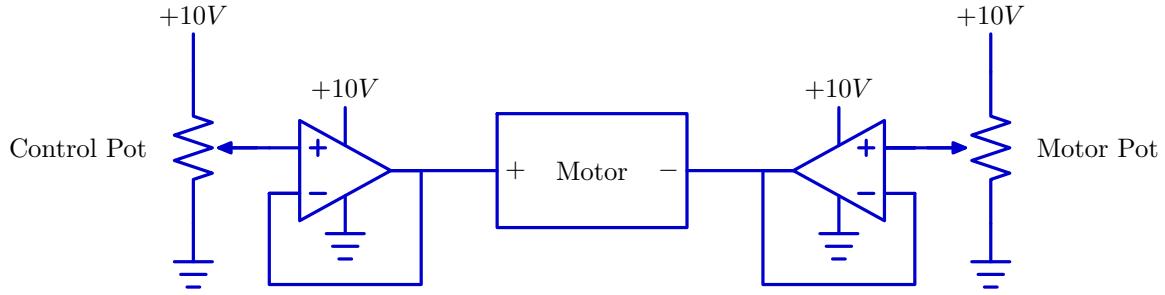


Figure 2-2: Sample circuit schematic: head angular position controller.

2.1.3 Protoboard

Protoboards are boards on which one can quickly build and test small circuits. They present a 2-dimensional array of interconnected dots in which circuit pieces and wires can be inserted. Figure 2-3 presents an example of an empty physical protoboard. In the orientation depicted in Figure 2-3, a protoboard has 4 groups of rows: the first 2 rows, the next 5 rows, the next 5 rows, and finally the last 2 rows. In the first and last groups, the dots on the protoboard are interconnected horizontally. In the middle two groups, the dots are interconnected vertically. This interconnection scheme is depicted in Figure 2-3.

2.1.4 Protoboard Layout

The protoboard layout of a given schematic is the placement of circuit pieces and wires on a protoboard that corresponds to the schematic. This is done by placing the appropriate pieces on the protoboard and then appropriately interconnecting them with wires as prescribed by the schematic. As an example, Figure 2-4 presents one possible protoboard layout corresponding to the schematic shown in Figure 2-2.

For each of the circuit components we are interested in, there is a corresponding circuit piece that may be inserted into the protoboard. The one exception is that op-amps come in pairs. That is, each op-amp circuit piece that is inserted in the protoboard actually contains two op-amp components within it. This raises an important design problem when laying out a schematic - the need to find a good way

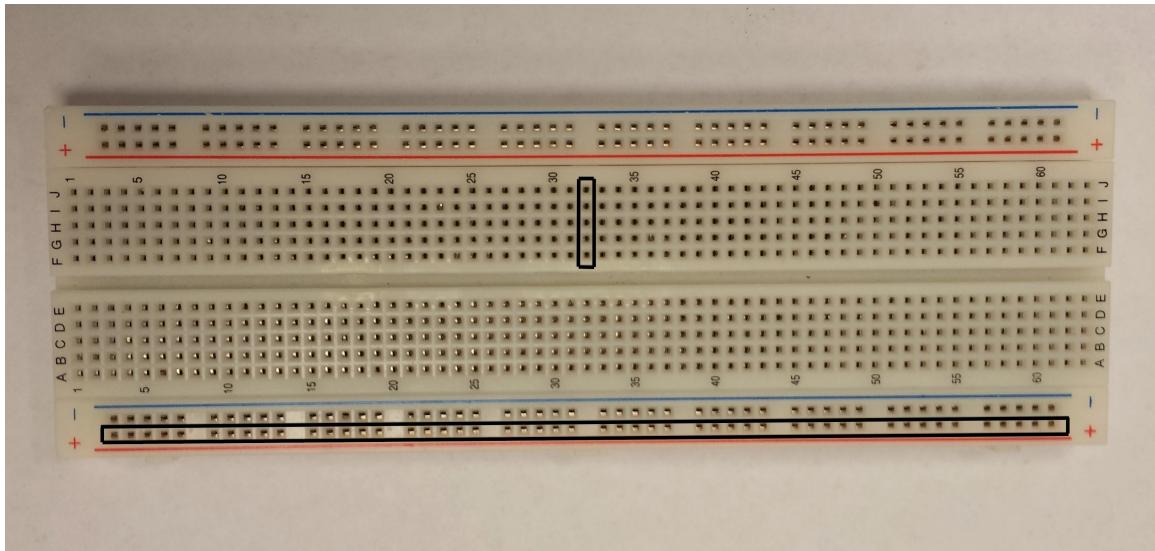


Figure 2-3: Physical protoboard. In the top and bottom groups, the dots are interconnected horizontally. In the middle two groups, the dots are interconnected vertically.

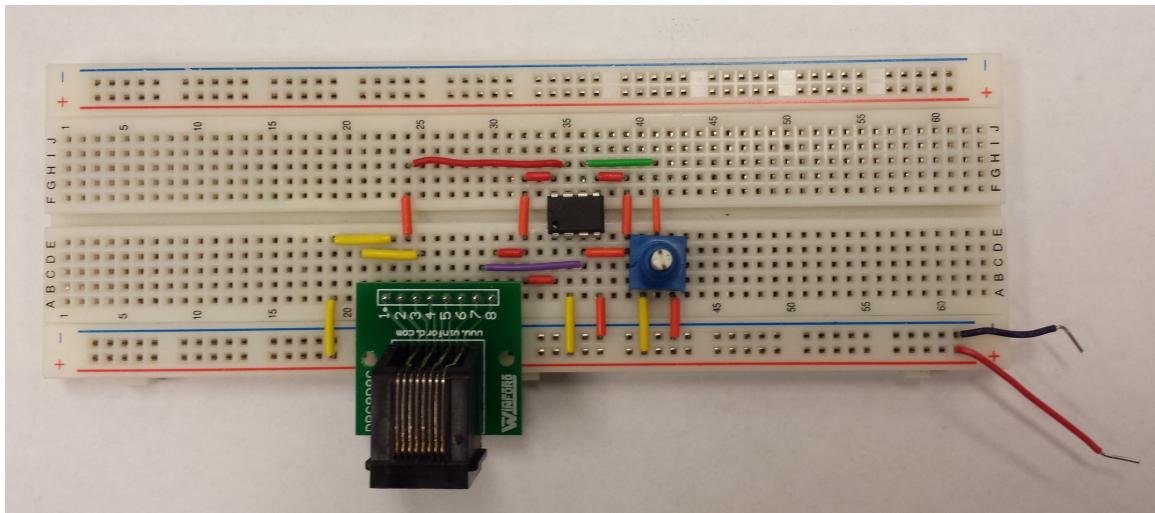


Figure 2-4: Protoboard layout for the schematic in Figure 2-2.

to group together the op-amps in a schematic to result in a “good” layout. To solve this problem, the designer must have some criteria for what makes a layout “good.” While there are no conclusive answers to this question, keeping in mind that we want layouts that are easy to build, easy to debug, and aesthetically pleasing, we could come up with the following rules of thumb:

- The layout should not have any wires that cross circuit pieces.
- The layout should have no crossing wires, especially occlusions (i.e. crossing wires with the same orientation).
- The layout should only have horizontal and vertical wires (i.e. no diagonal wires).
- The layout should have as few wires as possible.
- The total length of wires in the layout should be as small as possible.

Given the background information discussed thus far, the goal of this project is automatically generating a “good” protoboard layout from a circuit schematics.

2.2 Previous Work

Here we will discuss previous work that has been done relating to this project. First, as our project aims to augment the quality of 6.01, we look at the currently available infrastructure. Next, we look at what work has been done relating to layout in general.

2.2.1 CMax

In a typical circuits lab in 6.01, students first design their circuits by drawing schematics of the circuits on paper. After iteratively improving their designs based on discussions with staff members, they build their circuits on a simulation tool called Circuits Maximus (CMax)[2]. With CMax, a student can lay out a circuit on a simulated

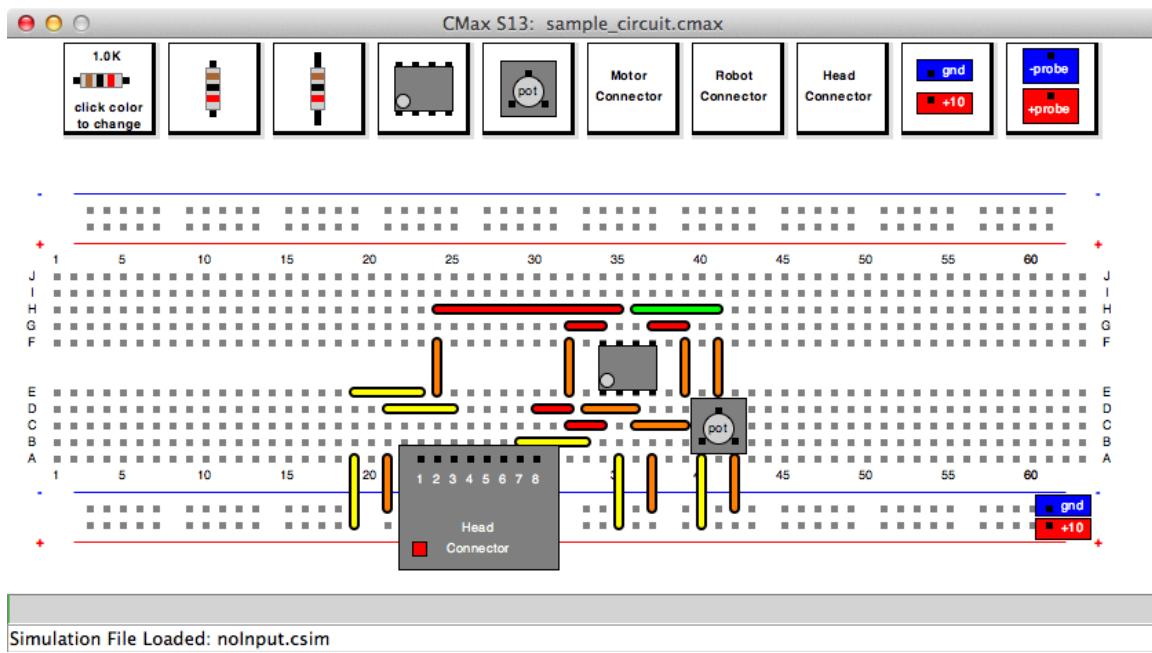


Figure 2-5: CMax layout for the schematic in Figure 2-2.

protoboard, and test the circuit to make sure that it behaves as desired. CMax provides a very fast and safe way of debugging circuit layouts compared to debugging layouts on a physical protoboard. Once the students are satisfied with their observations from CMax, they build their circuits on physical protoboards and carry out the appropriate experiments. Figure 2-5 presents a CMax layout of the schematic in Figure 2-2.

Using CMax has reduced circuit debugging time for 6.01 students. Its introduction has made learning circuits significantly easier for many students, especially those that have little or no prior experience with circuits. In addition to making the lab exercises more manageable, it provides students with a handy way to build, analyze, and experiment with circuits at their own leisure outside of lab.

While CMax is a fantastic tool, a tool that can automate the layout process can be even more useful. The most instructive part of the labs that students do in the circuits module of 6.01 is really designing the circuits in the first place, which they currently do by drawing schematic diagrams on paper. Once they are happy with their schematic diagrams, they proceed to laying out the corresponding circuits with CMax.

The process of laying out a schematic does not really have very much instructive substance. This process is essentially solving a puzzle, and has almost nothing to do with the subject matter – designing circuits. In fact, when the circuits get complicated and involve many pieces, translating a schematic diagram into a protoboard layout gets to be quite challenging and time-consuming. In these situations, students often end up with convoluted and unpleasant layouts that are very difficult to debug in the likely case of the circuit not behaving as expected. Not only are such layouts difficult to debug for the students themselves, but they are also often difficult for staff members to understand. In the best case scenario, students should have to work out the right schematic diagram of the circuit of interest, but should not have to produce the corresponding protoboard layouts.

With the schematic entry tool this paper introduces, a typical 6.01 circuits lab would proceed as follows. First, as before, the students would draw schematic diagrams of their circuits on paper. Once they have schematic drawings they are happy with, they can directly draw their schematic drawings on the schematic entry tool. In fact, students may proceed directly to building the schematic drawings on the tool, bypassing the experimentation on paper. Once they have a schematic drawn, they can analyze it with the tool, discuss it with staff members, and improve it easily and quickly with the tool. When they are satisfied with the behaviors of their schematic circuit, they can produce the corresponding protoboard layout automatically. The automatic generation of protoboard layouts would be the most important advantage of this tool. They can then build the layout on a physical protoboard and carry out experiments with it.

2.2.2 Current work in automatic layout

In my explorations, I was not able to find any tools that completely automatically translate circuit schematics into protoboard layouts. However, there do exist tools that perform partially- or fully-automatic Printed Circuit Board (PCB) layout. To my findings, most of these tools do not publish their algorithms and, rather, keep them proprietary. Hence, I was not able to build my work off of any existing products. In

a sense, this project aims to build something new.

Chapter 3

Methods

In this chapter, we discuss our solution to the problem stated in Chapter 1, as well as various alternatives we considered along the way. First, we briefly introduce the schematic entry GUI. Next, we discuss in detail how we solved the protoboard layout problem and how we evaluated our solution.

3.1 GUI

We designed the schematic entry GUI to have a rich set of features so as to make drawing schematics an easy and intuitive task for students. Figure 3-1 gives a version of the schematic shown in Figure 2-2 as drawn of the schematic entry tool. Appendix A discusses the features and capabilities of the schematic entry GUI in further detail.

3.2 Solving the Layout Problem

In broad terms, we solved the layout problem by formulating it as a search problem. Given a schematic of a circuit, we start from an empty protoboard, and search through the space of all possible protoboard layouts to find a protoboard layout corresponding to the schematic at hand. Importantly, we utilize various simplifications and heuristics to prune out many states in the search space.

We broke down the problem into two parts. The first task is finding a place-

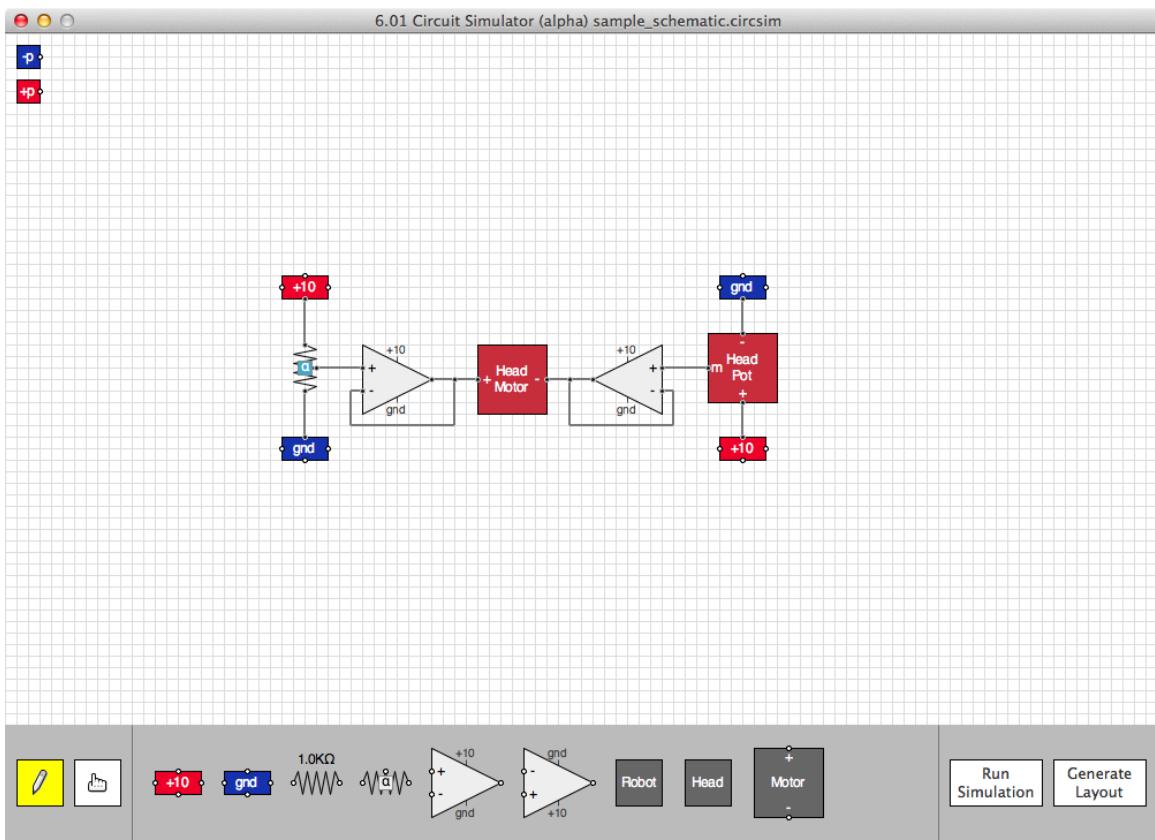


Figure 3-1: Sample schematic drawn on the schematic entry tool.

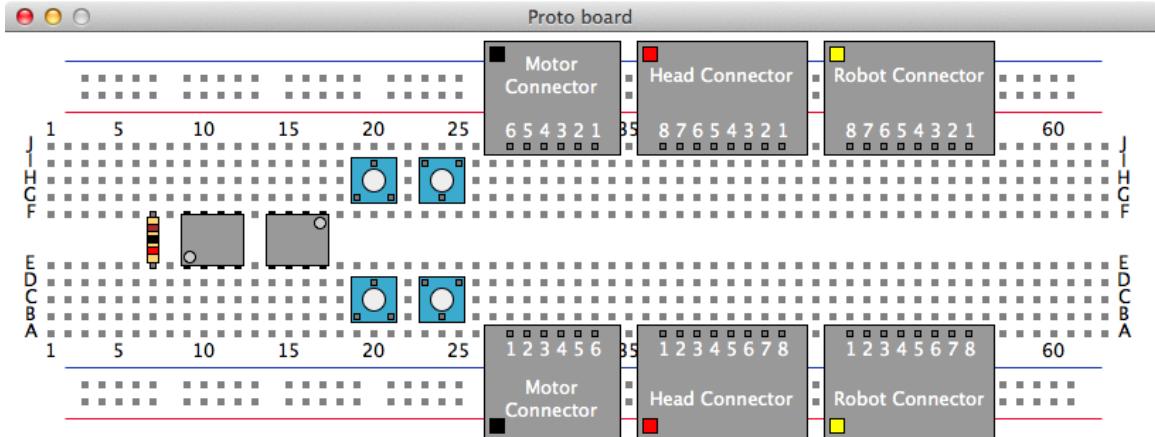


Figure 3-2: Various acceptable ways of putting each of the circuit pieces on the protoboard.

ment of all the circuit pieces on the protoboard. The second task is wiring them up appropriately.

3.2.1 Part 1: Placement

Let us first consider how to place a set of circuit pieces on the protoboard for a given circuit schematic. Any given circuit may contain resistors, op-amps, pots, motors, head connector parts, or robot connector parts. For each of these components, we must put down a corresponding piece on the protoboard. As each piece may be placed on the protoboard in one of many different ways, we first decided on a fixed set of allowed placements for each of the pieces. Figure 3-2 presents these acceptable placements. Resistors are placed in the middle strip of the protoboard. Pots have two possible vertical positions as well as two possible orientations. The connector pieces have two possible vertical positions each. Op-amp pieces are also placed in the middle strip of the protoboard, but with two possible orientations. Op-amp pieces are unique in that each op-amp pieces contains two op-amps within it. Thus, we face the task of packaging the op-amps in the schematic in the “best” possible way, i.e. so as to require as little work as possible when wiring the pieces together. Section 5.2 more precisely discusses the number of possibilities.

There are many available options when choosing a placement of a set of circuit pieces on the protoboard. First we must choose among a possibly large number of ways to package together the op-amps in the circuit. For each possible packaging of op-amps, we must consider various ways of placing the pieces on the protoboard, even with the restrictions put forth above.

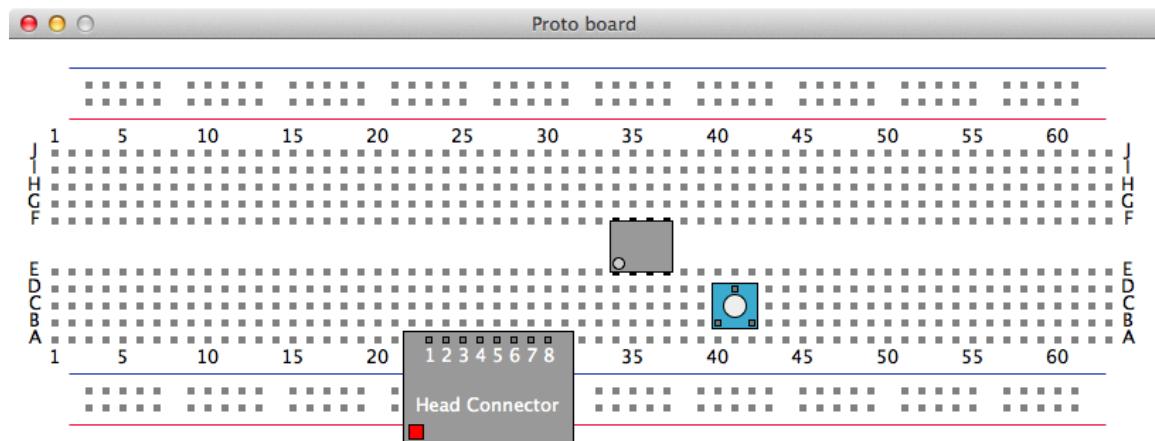
Simplifications

We reduce this large number of possibilities by only allowing placements in which no two pieces share a column. This is not necessary in general, but the number of pieces present in a typical 6.01 circuit could likely be realized under these circumstances.

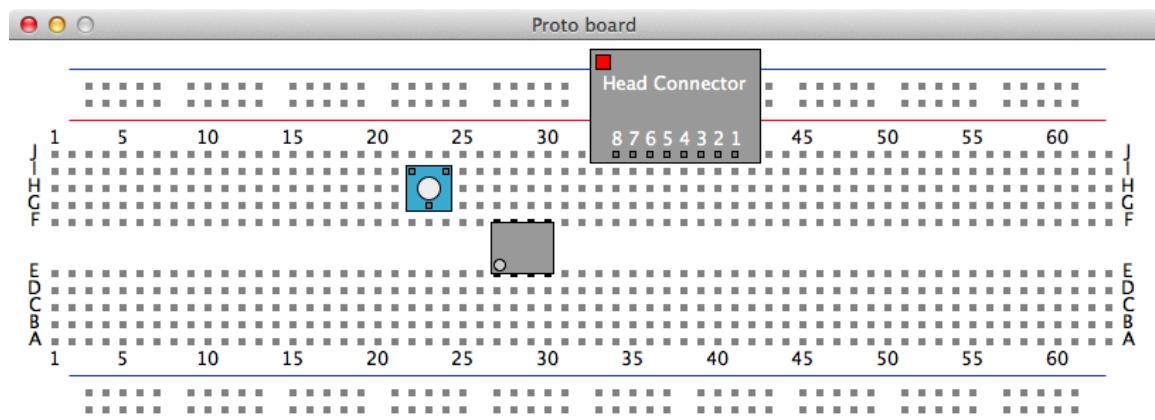
Next, we specify that there be exactly two columns on the protoboard separating each consecutive pair of pieces, unless the pieces are both resistors, in which case there must be exactly one column separating them. These numbers of columns were chosen to leave enough space for wiring. Given a set of pieces to be put on the protoboard, this specification reduces the problem of choosing a placement for the pieces to finding an *order* of the pieces together with choosing their respective vertical locations and orientations. Figure 3-3 shows two alternative placements for the schematic shown in Figure 3-1 that both respect the the stated placement specifications. We consider a few alternatives to finding a placement under these specifications.

Random Placement

One simple placement strategy is to choose a placement randomly. That is, we choose an op-amp packaging randomly; we choose an order of the pieces randomly; and we choose the vertical locations and orientations of the pieces randomly as well. The advantage of this approach is that it gives us a placement very quickly without requiring much computation. On the other hand, we may end up placing two pieces that need to be connected to each other very far apart, which would make the wiring task more difficult. We ought to consider alternatives in which we take into account the task of wiring. We should try to place the pieces so as to require as little work as possible during the wiring step.



(a) Placement 1



(b) Placement 2

Figure 3-3: Two possible placements for the schematic shown in Figure 3-1.

Minimal Heuristic Cost

The key idea is that if two pieces are meant to be connected together by wires, then they should be placed close to each other on the protoboard. We can capture this idea by assigning heuristic costs to the placements and choosing a placement that produces the minimal heuristic cost. To that end, there are two heuristic cost functions we considered.

Distance Based Cost Given a circuit schematic and a corresponding placement of the circuit pieces on the protoboard, every pair of components in the schematic that is connected by a wire gives us a corresponding pair of locations on the protoboard that must be connected by wires. However, we can express this requirement a little bit more concisely. We must consider all of the nodes in the schematic, and find the circuit components in the schematic that are connected to the respective nodes. Now for each node in the circuit, we get a set of locations on the protoboard that must be interconnected. The first step in devising the distance based cost function is to have a way to estimate the cost of connecting two locations on the protoboard. A simple such cost function is the Manhattan distance between the two locations. Since we want to produce layouts that only contain horizontal and vertical wires (i.e. no diagonal wires), the Manhattan distance cost is appropriate. Given this heuristic cost for connecting two locations with wires, we can define the cost for interconnecting the locations associated with a particular node to be the weight of the minimum spanning tree of the locations. Now we can define the cost of a placement to be the sum over all nodes in the circuit of the cost for interconnecting the locations for each node.

We can demonstrate this cost function using the two placements shown in Figure 3-4. In the Figure, each placement has two connections that must be made, the first indicated by two circled locations and the second indicated by two locations outlined by rectangles. The distance based cost for Placement 1 is $(3) + (2 + 3) = 8$ while the distance based cost for Placement 2 is $(7) + (2 + 3) = 12$. As desired, the cost function indicates that Placement 1 is the better of the two.

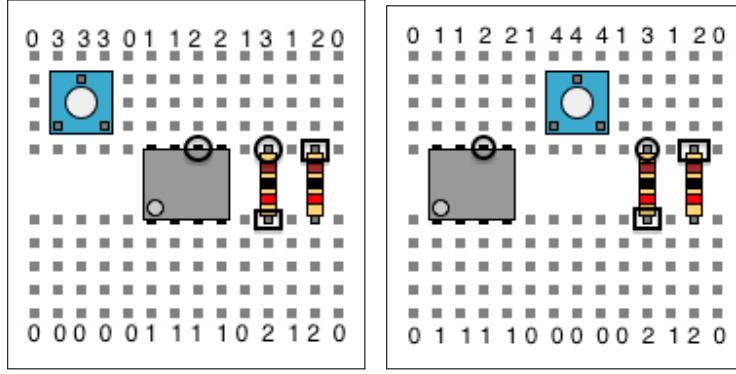


Figure 3-4: Placements to demonstrate cost functions.

Blocking Based Cost The most scarce resource on the protoboard are the rows. There are 10 rows with which to work, and in producing a layout, we must fit all the wires within these 10 rows. This gives us an idea on how to define how hard a placement will be to wire. Given a set of pieces, we can find a set of pairs of locations on the board that need to be connected as we did above. As noted in Section 2.1.3 there are two groups of 5 rows on a protoboard. For each group of rows and each column in the group of rows, we can count how many pieces reside in that column, and how many wires may pass through the column (in connecting a given pair of locations assuming an empty protoboard). The heuristic should reflect columns that would be used heavily. In our implementation, the final cost is computed as the sum of the squares of the counts for each column.

To demonstrate this cost function, let us look at the two placements in Figure 3-4 once again. Each column in the placements is labeled with the cost associated with the column in the manner described above. The cost for Placement 1 is the sum of the squares of the costs for each column, which evaluates to 65. The cost for Placement 2, computed similarly, is 87. Once again, this cost function indicates that Placement 1 is better.

Small Heuristic Cost

Using one of the two cost functions discussed above, we can aim to find a placement with the minimal cost. However, this involves trying all possible orderings of the pieces with which we are working. For example, if we are trying to order 10 pieces, we would need to look at $10! = 3628800$ possible orderings. Note that this is in addition to searching over all possible ways of packaging the op-amps together. It is clear to see that the search for a minimal cost placement quickly gets out of hand. Rather than looking for an optimal placement, we aim for a placement with small cost.

Algorithm 1 presents a polynomial-time procedure that orders a given list of pieces in a way that results in a small cost. The algorithm places one of the pieces at a time, starting from an empty placement. It relies on two ideas. First, once a piece has been placed, all the pieces that are connected to it will be placed soon after so that it is more likely that those pieces are placed close to it. Second, we place the pieces with the most nodes first since those are the ones that most likely have connections with many other pieces.

Algorithm 1: Producing a circuit piece placement with small heuristic cost.

Data: A list P of circuit pieces.

Result: A list R of circuit pieces representing a placement.

Sort P in decreasing number of nodes on the respective pieces

$Q \leftarrow$ empty Queue

$R \leftarrow$ empty List

while P is not empty **do**

Pop the first piece in P and push it onto Q

while Q is not empty **do**

$p \leftarrow Q.pop()$

Consider all vertical locations and orientations of p

Place p at an index in R that minimizes the cost of P

foreach piece q in P connected to p **do**

Pop q out of P and push it onto Q

3.2.2 Part 2: Wiring

Once the placement task is done, the next problem is wiring. We approach this problem as a search problem and use the A^* search algorithm to solve it. In fact, the wiring step uses an infrastructure for the A^* search algorithm exactly as presented in 6.01. Hence, students in the class may appreciate an application of something they learned earlier in the course to produce a tool that they are using for something that may seem completely unrelated and difficult.

Using A^*

The A^* algorithm can be used to search for a path from some starting vertex¹ in a graph to some goal vertex². The algorithm works by keeping track of an agenda of vertices to consider in a priority queue, where the value associated with each vertex in the priority queue is the sum of the cost to get from the start vertex to the vertex at hand and the value of the heuristic computed at the state of the vertex, which is an estimate of the minimum cost to get from the vertex at hand to a goal vertex. At each step, the algorithm pops one vertex from the priority queue (the vertex with the minimum associated value). If the vertex happens to satisfy the goal of the search, the algorithm returns the state for that vertex as the answer to the search problem. Otherwise, it adds the children vertices of that vertex to the priority queue and continues. When adding children vertices, the algorithm takes care not to reconsider states that it has already considered via a different path. We call the process of popping a vertex from the priority queue and treating it as described *expanding* the vertex. In general, when using the A^* algorithm, we need to design four things:

1. The notion of a vertex in the search tree, the cost associated with a vertex, and how we obtain the neighbors of a vertex,

¹The preferred terminology is “a node in a graph” but here we will use the term “vertex” since we already use “node” to refer to nodes in circuits.

²In fact, A^* guarantees an optimal path, a path that has the minimum possible cost from the starting vertex to a goal vertex, if we use a heuristic that is *admissible*. A heuristic is said to be admissible if it does not overestimate the actual minimal cost to a goal vertex for any state. Here, however, we will not worry about the admissibility of our heuristic as our main goal is designing the heuristic to prune as many states as possible.

2. The starting vertex,
3. How we identify whether a particular vertex in the search tree achieves the goal of the search, and
4. A heuristic function that estimates the distance from a given vertex to a goal vertex.

Vertices

Each graph vertex will represent a protoboard layout and a set of locations on the protoboard that have yet to be connected by wires. The starting vertex will represent a protoboard layout that just has the circuit pieces with no wires as well as all the pairs of locations that must be connected.

We obtain the neighbors of a vertex by taking the current protoboard layout and producing new ones in which we place exactly one wire at various locations. We choose the starting point of a wire to be any one of the free locations on the protoboard that is already connected to one of the pieces, and we extend the wires in all possible vertical and horizontal directions up to some fixed wire length. Note that we need to take great care when placing wires in order not to short, or directly connect, two different nodes. We discard any vertices that arise from placing a wire that shorts two different nodes.

The way we define the cost of a vertex, i.e. the cost of getting from the starting vertex to a vertex of interest, depends on our definition of a good protoboard layout. In general, we want to penalize having long wires, many wires, or crossing wires. In our implementation, while we have a large penalty for two crossing wires of opposite orientations (i.e. vertical and horizontal), we do not allow occlusions as they are particularly difficult to physically build and debug. Finally, we want to favor making a desired interconnection between locations on the protoboard. That is, if placing one wire results in a layout in which one of the pairs of locations that need to be connected becomes connected, then the cost of that child vertex should reflect that fact.

An important consideration we need to make is how to organize the search. Recall that we have a set of nodes in the circuit of interest, and for each node we have a set of locations that need to be interconnected. We considered the following six different strategies to carry out the search:

1. All pairs: For each node, collect a set of pairs of locations on the protoboard corresponding to a minimum spanning tree of the locations for that node, so that if all pairs of locations in this spanning tree are connected, then the locations for the node will be interconnected. Collect all such pairs of locations for all of the nodes in the circuit, and have the starting vertex represent this set of pairs of locations. In this strategy, we run exactly 1 search to solve the problem.
2. Per-node (increasing): Treat each node separately. That is, iteratively connect the locations for each of the nodes until there are no more disconnected nodes in the circuit. In this strategy, we run a number of searches equal to the number of nodes in the circuit. Order the searches in *increasing* order of the number of locations per node, breaking ties arbitrarily.
3. Per-node (decreasing): Similar to per-node (increasing), but order the searches in *decreasing* order of the number of locations per node.
4. Per-pair (increasing): Treat each pair of locations that needs to be connected separately. That is, iteratively connect pairs of locations that need to be connected until there are no more disconnected pairs. In this strategy, we run a number of searches equal to the number of pairs of locations that must be connected. Order the searches in *increasing* order of the Manhattan distance between the pairs of locations, breaking ties arbitrarily.
5. Per-pair (decreasing): Similar to per-pair (increasing), but order the searches in *decreasing* order of Manhattan distance between the pairs of locations.
6. Straight: As a back-up alternative, we consider using one (possibly diagonal) wire to connect each of the pairs of locations that must be connected. This approach requires no search and does not take layout quality into consideration.

The strategy we choose among these six has a significant effect on the outcome of the wiring step. We will look at the difference in detail in Chapter 4.

Goal test

We test that a given vertex is a goal vertex by verifying that its representation indicates no further pairs of locations to connect.

Search heuristic

In A^* search, choosing the right heuristic can often make the search much more efficient. Given a vertex, we can estimate its distance from a goal as follows. For each pair of locations $(\text{loc}_1, \text{loc}_2)$ that need to be connected, we could consider the pair's distance from the goal to be the smallest Manhattan distance between any location connected to loc_1 and any other location connected to loc_2 . To compute the heuristic cost of a vertex, we simply add up this value for each of the pairs of locations that need to be connected. In Chapter 4 we will compare the performance of A^* with this heuristic versus carrying out Best First Search with this heuristic. In Best First Search, as opposed to in A^* , vertices are considered in order of decreasing heuristic value, without consideration for the cost incurred on the path from the starting vertex.

Limiting the number of expanded vertices

With the implementation of A^* discussed so far, the algorithm terminates if we either find a solution, or we exhaust the search space without finding a solution. In our search problem, the number of states is very large (Section 5.1 discusses the search space size in more detail), so this implementation of A^* may sometimes run out of memory during the search. To mitigate this problem, we introduce a limit to the number of vertices the algorithm expands before giving up. That is, if the algorithm expands a certain fixed number of vertices and still has not found an answer, the algorithm gives up. For us, this limit is set to 300 vertices. In Chapter 4 we provide data that describes the effect of this choice.

3.2.3 Combining the Methods

With the methods discussed so far, we aimed to completely solve the layout problem with one placement method and one wiring method. However, as we will soon see, such an algorithm is bound to fail on some set of schematics. When we ultimately put the final algorithm in front of students, we would like to avoid failure. The algorithm should be able to generate a layout for any schematic. Generating a layout with a few diagonal or crossing wires is better than silently failing and leaving the student empty handed. Here, we will discuss how we combine the methods described so far into one layout algorithm. The motivation for this combination will be discussed in Chapter 5 based on the data we obtain for the alternatives described thus far. Algorithm 2 presents the combined algorithm.

Algorithm 2: Layout algorithm obtained by combining multiple alternatives.

Data: A circuit schematic C .

Result: A protoboard layout corresponding to C .

```
foreach Placement cost metric  $M$  in (DISTANCE, BLOCKING) do
     $P \leftarrow$  Placement for  $C$  by using cost metric  $M$ .
    Connect the top and bottom rails on  $P$ .
    foreach Order  $O$  in (INCREASING, DECREASING) do
         $pairs \leftarrow$  Pairs of location on  $P$  to connect given schematic  $C$  and
        connection order  $O$ .
        foreach  $(loc_1, loc_2)$  in  $pairs$  do
            Attempt to connect  $loc_1$  and  $loc_2$  on  $P$ .
            If successful, update  $P$  accordingly and then post-process  $P$ .
            If not, record that the pair  $(loc_1, loc_2)$  was not successfully
            connected.
        if all pairs are successfully connected, return  $P$ .
    Pick unfinished layout with fewest and most compact disconnected pairs.
    Connect remaining pairs with shortest possible wires (possibly diagonal).
    Post-process and return resulting layout.
```

There are a few key items in Algorithm 2. The algorithm works by attempting to solve the problem in four different ways: two different ways of doing placement together with two different orders of wiring pairs. Note, therefore, that we have chosen the per-pair wiring scheme discussed above. If any one of the four trials succeeds, the algorithm immediately returns the corresponding layout. If all four trials fail,

on the other hand, the algorithm picks one of the four unfinished layouts that would require comparatively little additional wiring and completes that solution by placing possibly diagonal wires. This last step makes it highly unlikely that the algorithm will ever fail. The only way for the algorithm to fail is for there to be two nodes on the protoboard that need to be connected where all of the protoboard locations for at least one of the nodes are occupied, which is highly unlikely. This high success rate comes at the cost of placing wires that will almost surely cross other wires on the board.

The algorithm automatically starts out by connecting the top and bottom rails of the protoboard so that all rail rows are used to connect to power and ground, and no other nodes. This is a restriction that makes it easier to debug and amend the resulting layout.

The algorithm also has a post-processing step that attempts to improve the layout. The post-processing step makes three types of simple changes to the layout. First, we throw away any superfluous wires that do not serve to connect two parts of the circuit. Superfluous wires may be added to the layout in the search done by the wiring step, though very rarely. Second, we truncate long vertical wires into an equivalent set of smaller wires. For example, a wire going from one of the top rails to one of the bottom rails can be replaced by three smaller wires making the same connection. This change frees up rows for subsequent connections. Finally, if shifting a horizontal wire up or down results in a layout with fewer crossing wires, we make that change.

The last important aspect of this final algorithm not explicitly stated in Algorithm 2 is that the algorithm that will be put in front of students will only be allowed to use wires of a select few lengths. The kits that students work with do not come with wires of all lengths, so we force the wiring step to use wires of only those allowed lengths. We also avoid using length-1 wires because they are difficult to insert and remove from physical protoboards and are also difficult to see.

3.2.4 Evaluation

Here we present how we evaluated our solution to the layout problem to test how well it would serve students in 6.01. We tested the layout tool on numerous schematics and analyzed its performance on generating layouts from those schematics. As manually generating numerous test schematics is tedious and time-consuming, we devised a method to randomly generate thousands of test schematics. As the tool targets 6.01 labs, we tried to design the randomly generated schematics so that the range of complexity of these schematics mimics the range of complexity of circuits that students may build in 6.01.

The random schematic generation goes as follows. We create 6 basic parts of schematics. These 6 bases are:

- Three resistors arranged in a T-shaped configuration.
- Two resistors in series connected to a follower op-amp configuration.
- A pot connected to a follower op-amp configuration.
- A motor.
- A robot head.
- A robot.

These bases are depicted in Figure 3-5. They cover all of the components that may be necessary in a 6.01 circuit. Each base offers at least 3 points of connection with other bases. The random generation algorithm takes all possible combinations of up to 6 bases, allowing for repetition of bases with some restrictions. The robot head and robot bases can appear at most once as there is no need for more than one of each of these in 6.01 labs. The pot and follower op-amp base can appear at most twice as we never need more than two pots in 6.01 circuits. The motor base can also appear at most twice as we never need more than two motors per circuit in 6.01 labs. The other two bases, T-resistor configuration and two resistors in series with a follower op-amp, can be repeated up to 6 times. For a given combination of bases, we generate a set

of schematics in which we randomly choose to make connections between the bases.

Figure 3-6 presents a sample randomly generated schematic.

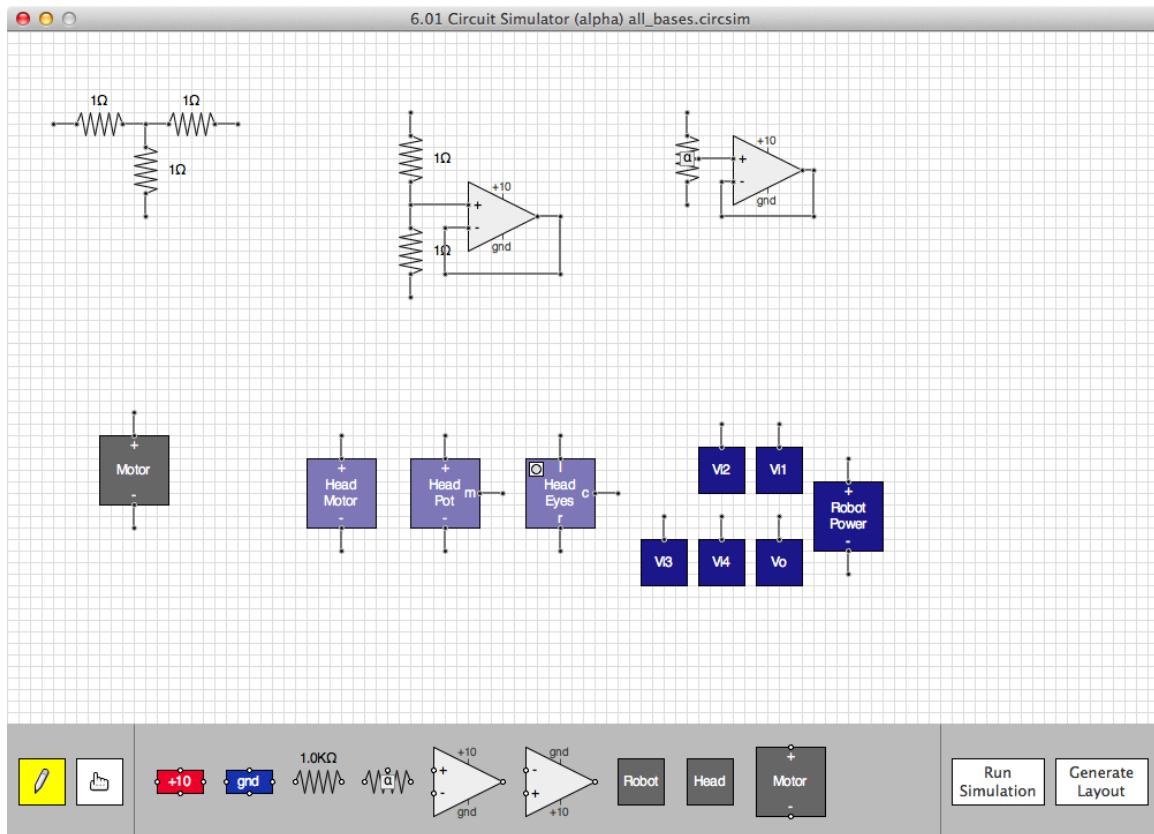


Figure 3-5: Bases for random schematic generation.

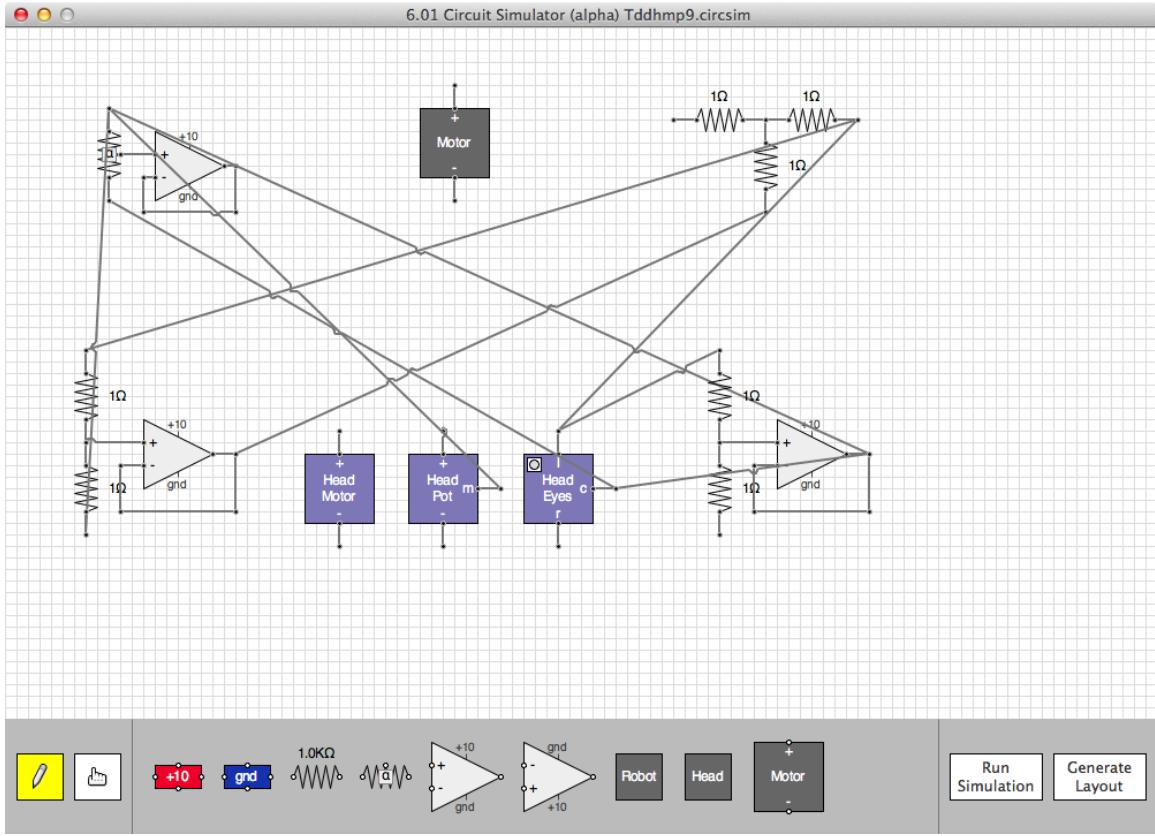


Figure 3-6: Sample randomly generated schematic.

This scheme produces a total of 4425 test schematics out of $1.2 \cdot 10^{27}$ possibilities³. When testing a particular algorithm on these test schematics, we run the algorithm on each test schematic 10 times. Chapter 4 presents the data collected in this manner comparing the various alternatives discussed in this Chapter.

An important question we must answer is how we quantify the goodness (or badness) of a particular layout. Our approach takes a weighted sum of a particular set

³These numbers were computed computationally as the variety of choices that could be made in the process makes it difficult to compute them combinatorially.

of attributes of a given layout. We define the badness of a layout to be:

$$\begin{aligned}
 & 1 \cdot \text{NUMBER_OF_WIRES} & + \\
 & 2 \cdot \text{TOTAL_WIRE_LENGTH} & + \\
 & 10 \cdot \text{NUMBER_OF_WIRE_CROSSES} & + \\
 & 10 \cdot \text{NUMBER_OF_DIAGONAL_WIRES} & + \\
 & 50 \cdot \text{NUMBER_OF_WIRE-PIECE_CROSSINGS} & + \\
 & 500 \cdot \text{NUMBER_OF_WIRE_OCCLUSIONS}.
 \end{aligned}$$

We will use this metric to decide which of a given set of alternative layout generation strategies tends to produce better layouts. The weights in the metric were chosen to reflect how bad each of the attributes are relative to one another. This choice of weights, therefore, reflects the following reasonable set of statements. Recall that our goal is to produce layouts that are easy to build, easy to debug, and aesthetically pleasing.

- Having an additional wire is about as bad as increasing the total wire length on the protoboard by 2.
- Having two wires that cross is about as bad as increasing the total wire length on the protoboard by 5.
- A diagonal wire is about as bad as two wires that cross.
- Having 10 wires that cross is about as bad as having a wire that crosses a circuit piece.
- Having 10 wires that cross circuit pieces is about as bad as having a wire occlusion. The penalty on wire occlusions is meant to indicate that we want absolutely no wire occlusions in our layouts.

Chapter 4

Results

In Chapter 3, we discussed a general solution to the protoboard layout problem and various alternatives that can be used in implementing the solution. Figure 4 summarizes the alternatives. In this chapter, we explore these alternatives and compare them quantitatively. We provide useful data for comparing the alternative strategies, and the data will be discussed in detail in Chapter 5.

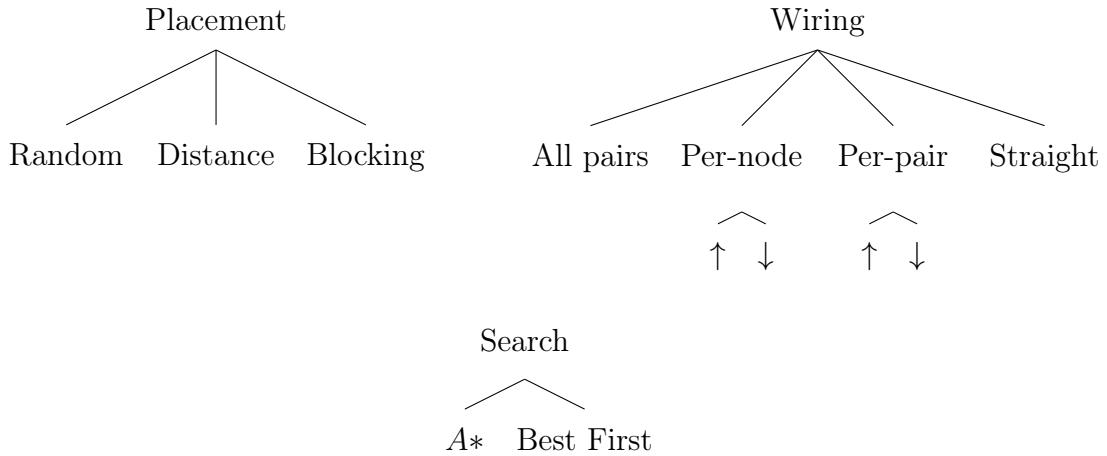


Figure 4-1: All possible alternatives to the algorithm.

As comparing all $3 \times 6 \times 2 = 36$ possible implementations of the algorithm is tedious, we will compare the alternatives for each aspect of the algorithm while holding other aspects fixed. Hence, we will carryout the following comparisons:

1. Placement: Distance vs. Blocking. Wiring method will be per-pair (decreasing),

and we will use A^* .

2. Wiring: All pairs vs. Per-node (decreasing) vs. Per-node (increasing) vs. Per-pair (decreasing) vs. Per-pair (increasing) vs. Straight. Placement method will be blocking, and we will use A^* .
3. Search: A^* vs. Best First. Placement method will be blocking, and wiring method will be per-pair (decreasing).

The data to compare the alternatives is gathered as described in Chapter 3. We run the algorithm on 4425 randomly generated schematics of varying complexities. The algorithm is run 10 times on each schematic.

In comparing alternatives, we consider 3 questions:

1. Which alternative is successful most often?
2. Which alternative, when successful, takes the least amount of time?
3. Which alternative, when successful, produces the best layouts?

We are also interested in how each of these attributes (success rate, running time, and layout quality) vary with circuit complexity. To quantitatively describe the complexity of a circuit, we look at the number of pins in the circuit, where a pin is defined to be a connection point on a circuit component that is connected by wires to another connection point (on the same component or a different component). Figure 4-2 presents a histogram of the number of pins in the schematics in dataset that will be used to do all comparisons in this chapter, not including the data presented in Section 4.5, for which we used a newly generated dataset of schematics to analyze the performance of the combined algorithm.




Figure 4-2: Histogram of the complexity the 4425 schematics used for evaluation.

To compare success rates, we will look at number of successes out of 10 runs on each of the 4425 schematics. To compare running time, we will look at CPU time spent on the wiring step, as the placement step has much less variability. To compare the goodness of layouts, we will look at numbers of wires, total lengths of wires, numbers of wire crosses, and our layout badness metric as functions of circuit complexity. Note that in all figures that follow, error bars indicate 1.96 times the standard error. For each comparison, we present exemplar layouts generated by the alternative methods for the schematic shown in Figure 4-3.

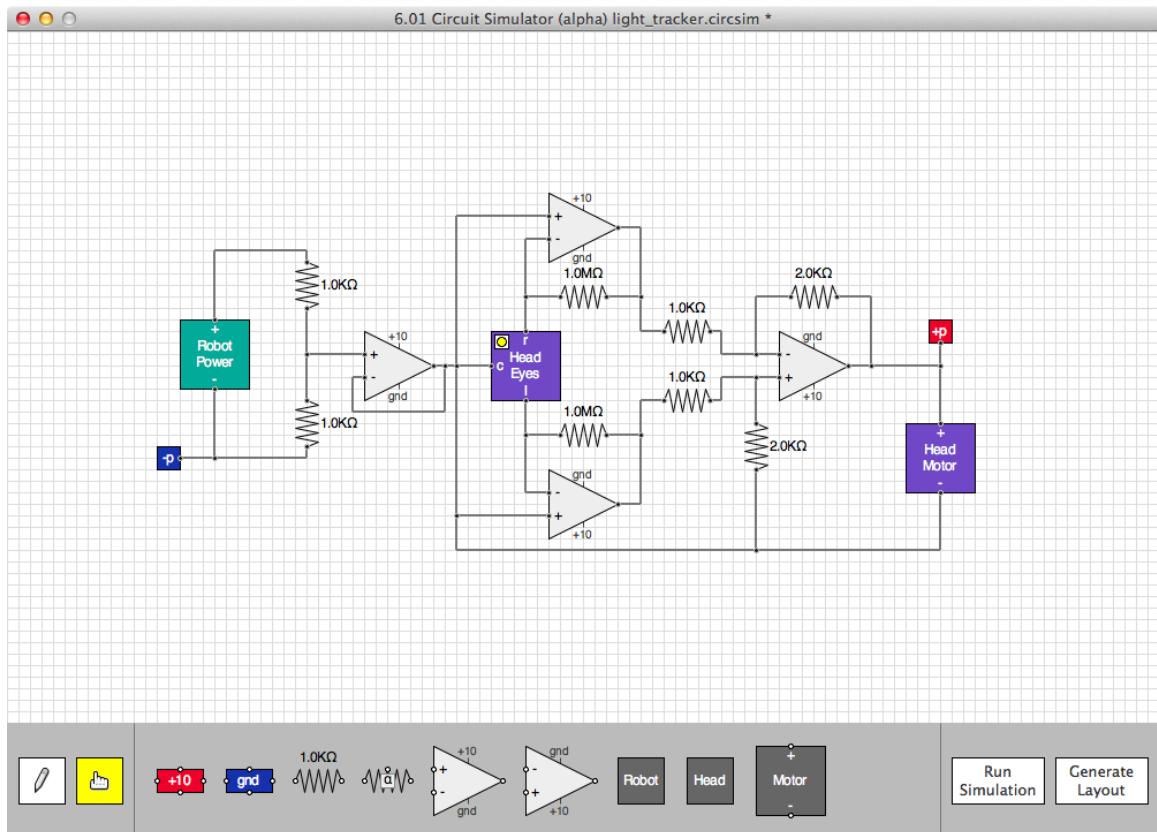


Figure 4-3: Exemplar schematic.

4.1 Comparing Placement Methods

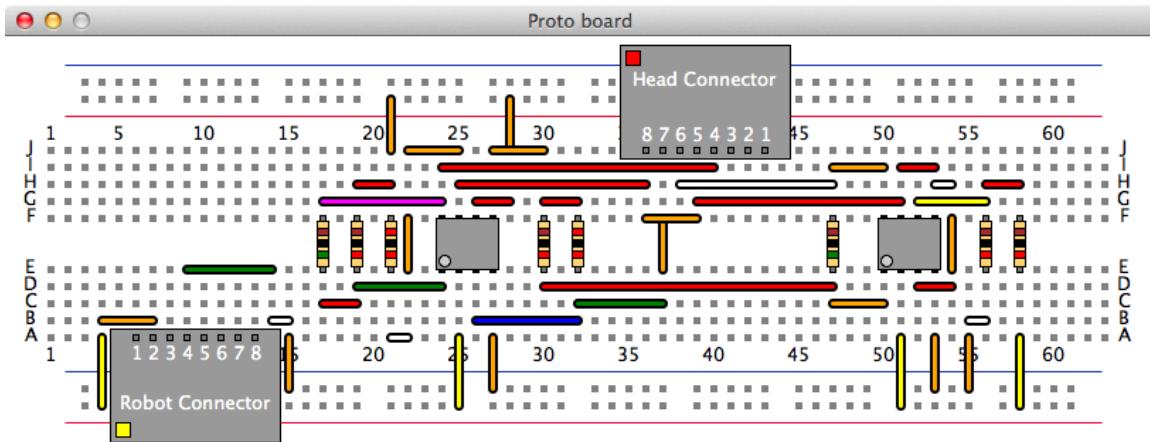


Figure 4-4: Blocking exemplar.




Figure 4-5: Distance exemplar.

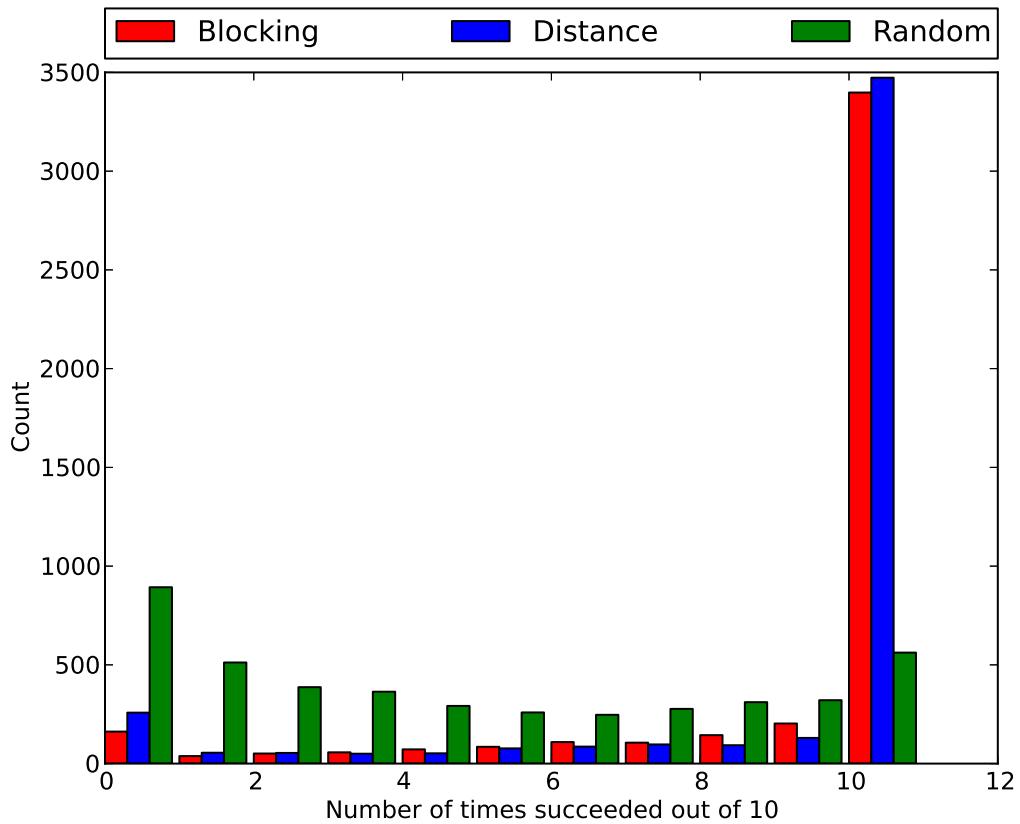


Figure 4-6: Placement method comparison: success rates.

	Number of times succeeded out of 10										
	0	1	2	3	4	5	6	7	8	9	10
Blocking	162 0.04	38 0.01	51 0.01	57 0.01	72 0.02	85 0.02	109 0.02	106 0.02	144 0.03	203 0.05	3398 0.77
Distance	258 0.06	55 0.01	54 0.01	50 0.01	52 0.01	77 0.02	86 0.02	97 0.02	93 0.02	130 0.03	3473 0.78
Random	893 0.20	512 0.12	387 0.09	364 0.08	292 0.07	259 0.06	247 0.06	277 0.06	311 0.07	321 0.07	562 0.13

Table 4.1: Placement method comparison: success rates.

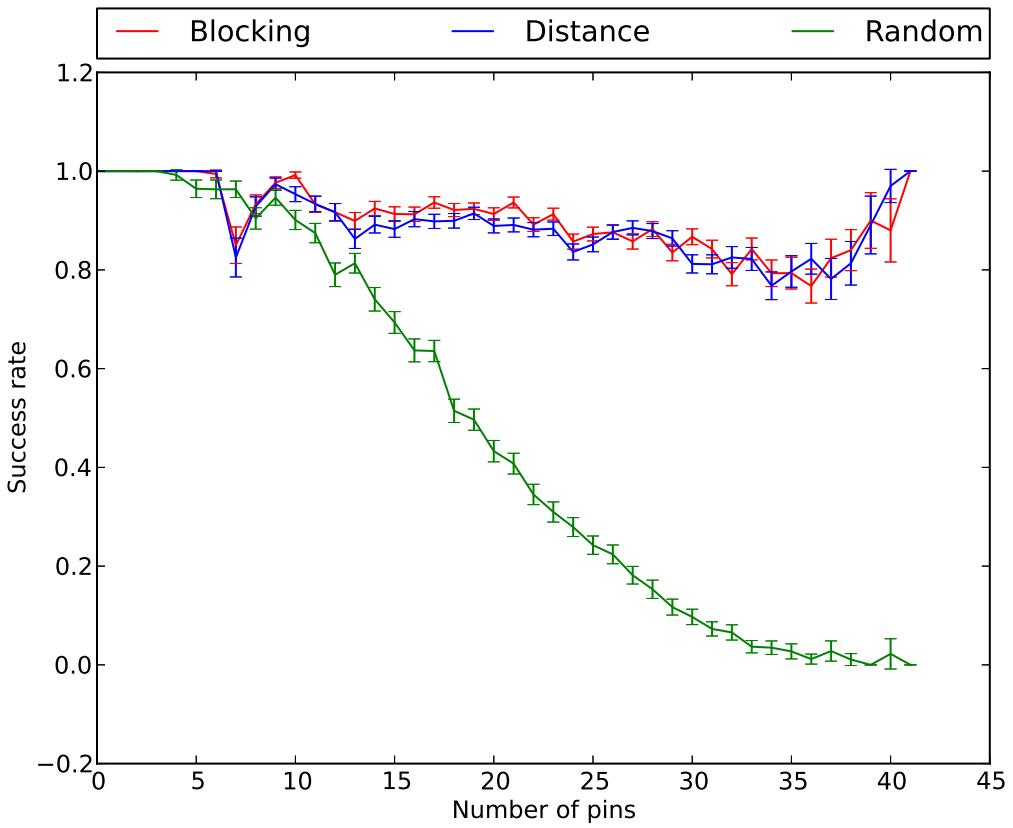


Figure 4-7: Placement method comparison: success rate trends.




Figure 4-8: Placement method comparison: wiring time trends.

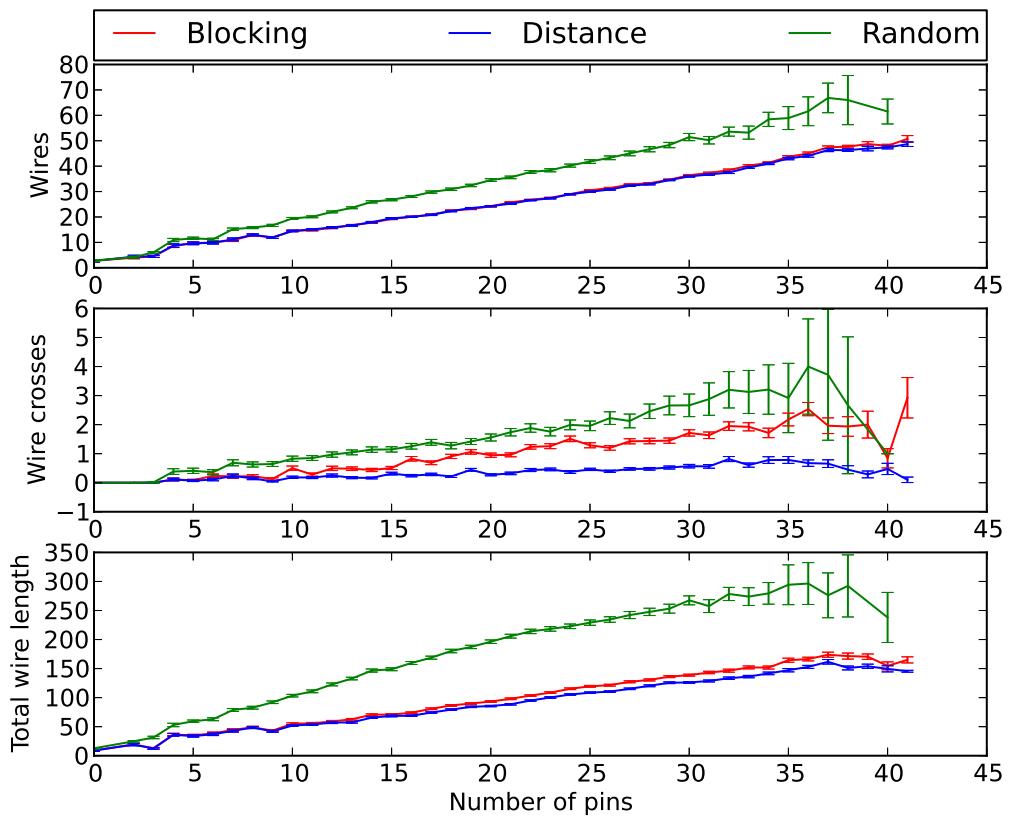


Figure 4-9: Placement method comparison: layout quality trends.

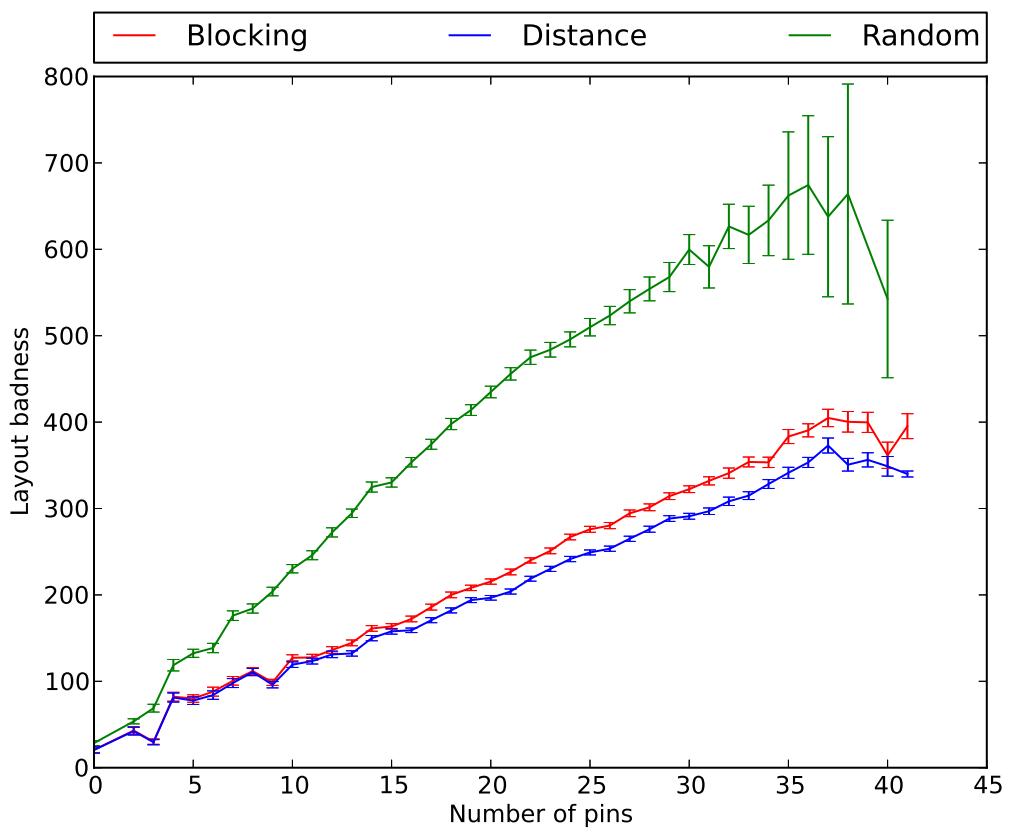


Figure 4-10: Placement method comparison: layout badness trends.

4.2 Comparing Wiring Methods

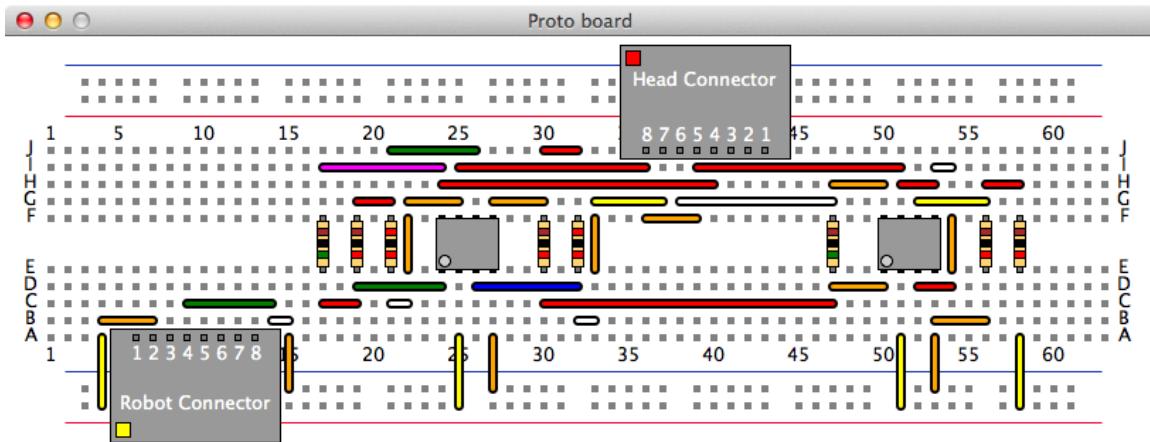


Figure 4-11: All pairs exemplar.

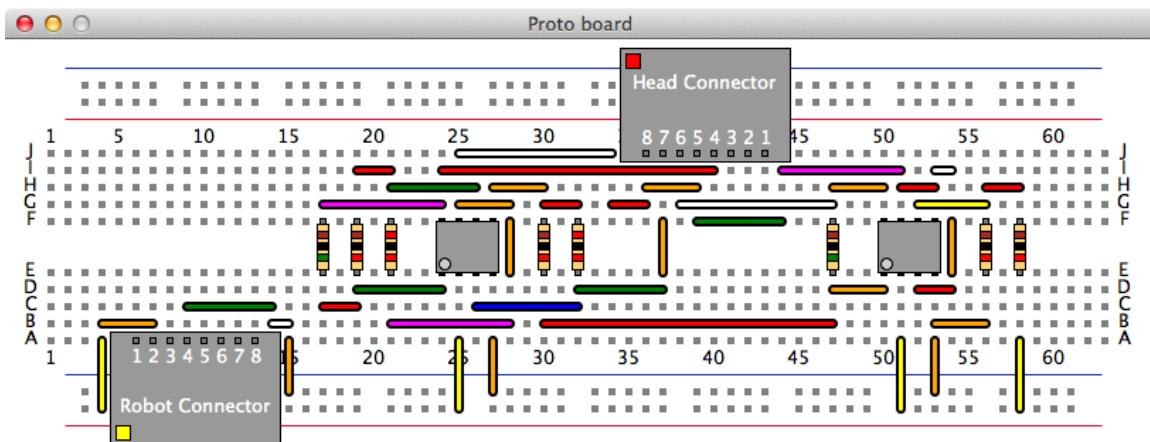


Figure 4-12: Per-node (increasing) exemplar.

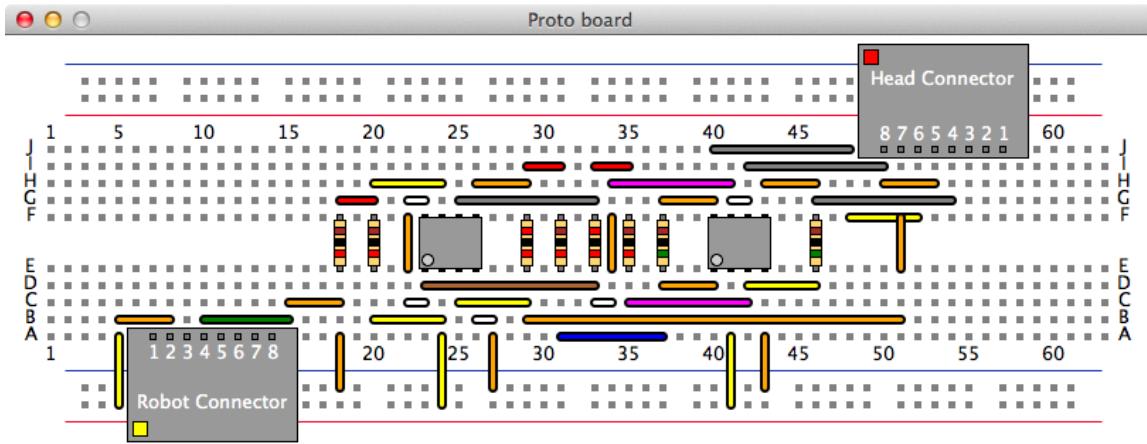


Figure 4-13: Per-node (decreasing) exemplar. We used the distance based placement method instead of the blocking based placement method to generate this exemplar as the combination of the blocking based method with this wiring method consistently failed on the schematic shown in Figure 4-3.




Figure 4-14: Per-pair (increasing) exemplar.

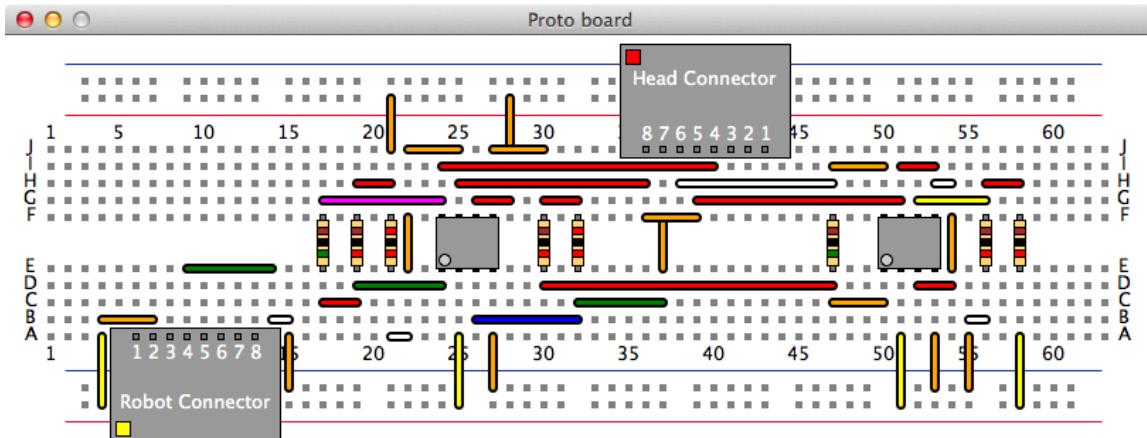


Figure 4-15: Per-pair (decreasing) exemplar.

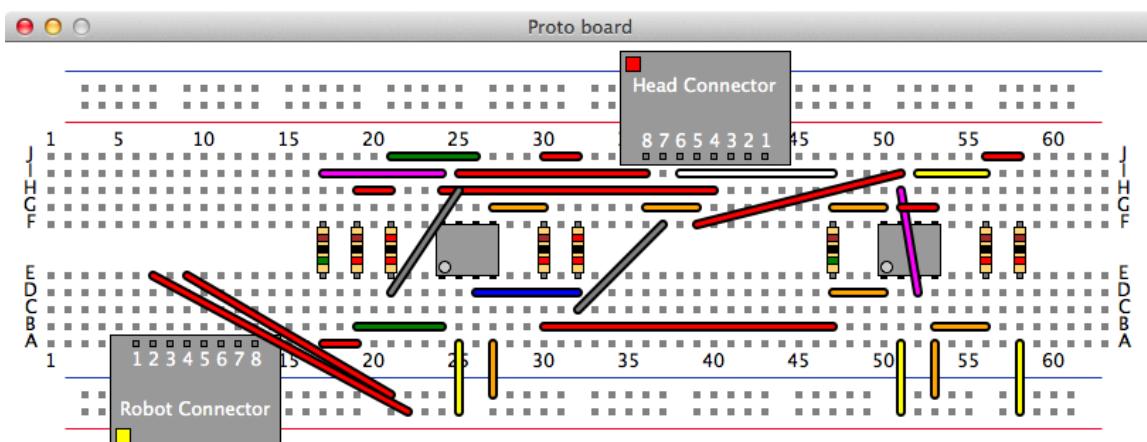


Figure 4-16: Straight wiring exemplar.

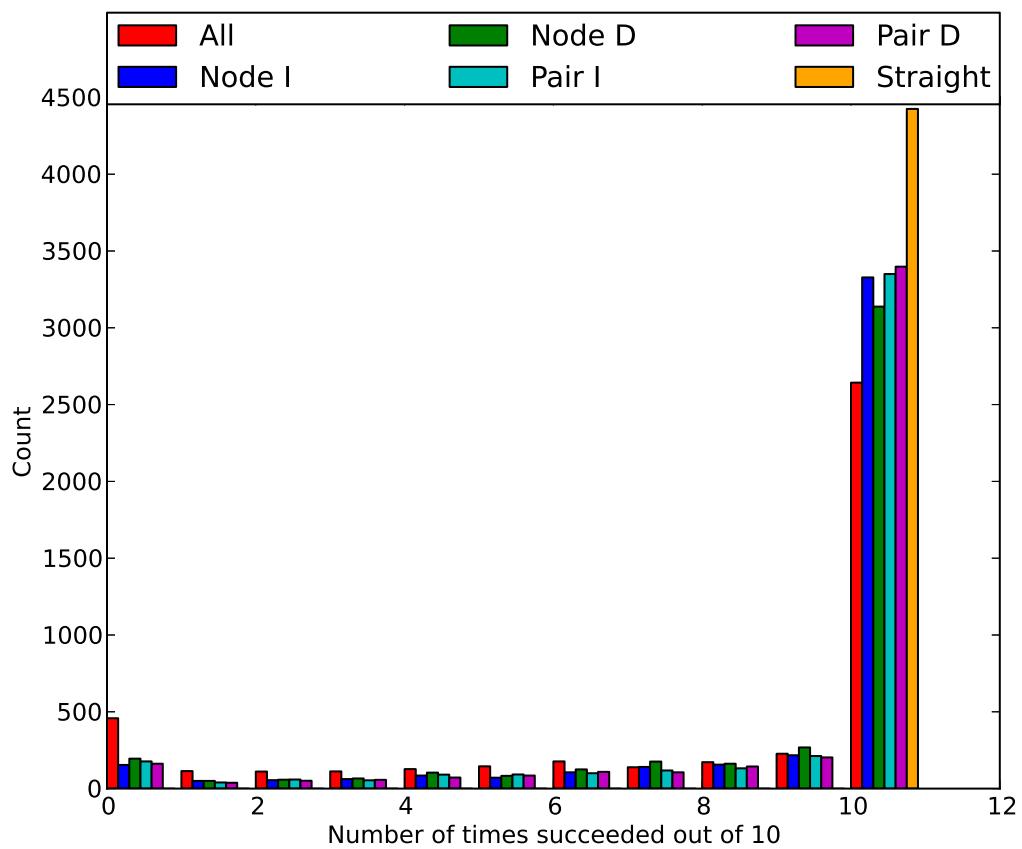


Figure 4-17: Wiring method comparison: success rates.

	Number of times succeeded out of 10										
	0	1	2	3	4	5	6	7	8	9	10
All	458 0.10	114 0.03	111 0.03	112 0.03	127 0.03	145 0.03	177 0.04	139 0.03	172 0.04	227 0.05	2643 0.60
Node I	154 0.03	50 0.01	55 0.01	62 0.01	85 0.02	71 0.02	106 0.02	141 0.03	156 0.04	217 0.05	3328 0.75
Node D	195 0.04	50 0.01	58 0.01	66 0.01	104 0.02	83 0.02	125 0.03	176 0.04	162 0.04	268 0.06	3138 0.71
Pair I	177 0.04	40 0.01	59 0.01	54 0.01	91 0.02	92 0.02	100 0.02	118 0.03	132 0.03	212 0.05	3350 0.76
Pair D	162 0.04	38 0.01	51 0.01	57 0.01	72 0.02	85 0.02	109 0.02	106 0.02	144 0.03	203 0.05	3398 0.77
Straight	0 0.00	0 0.00	0 0.00	0 0.00	0 0.00	0 0.00	0 0.00	0 0.00	0 0.00	0 0.00	4425 1.00

Table 4.2: Wiring method comparison: success rates.




Figure 4-18: Wiring method comparison: success rate trends.

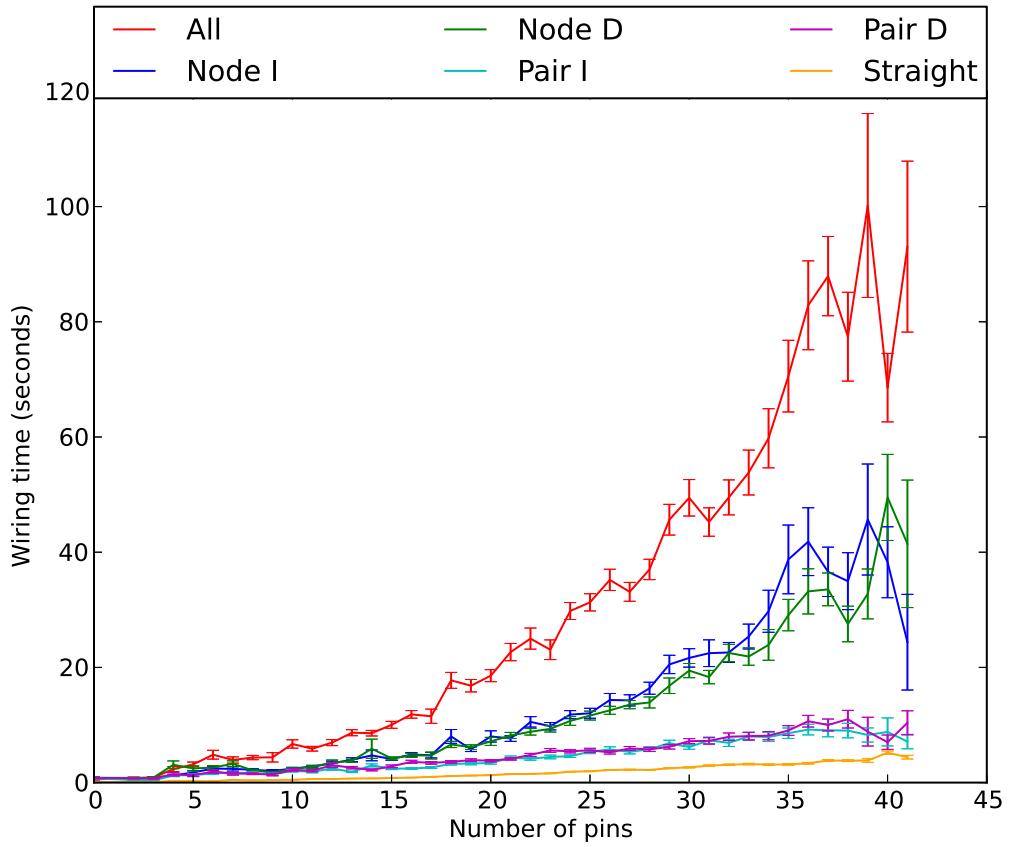


Figure 4-19: Wiring method comparison: wiring time trends.

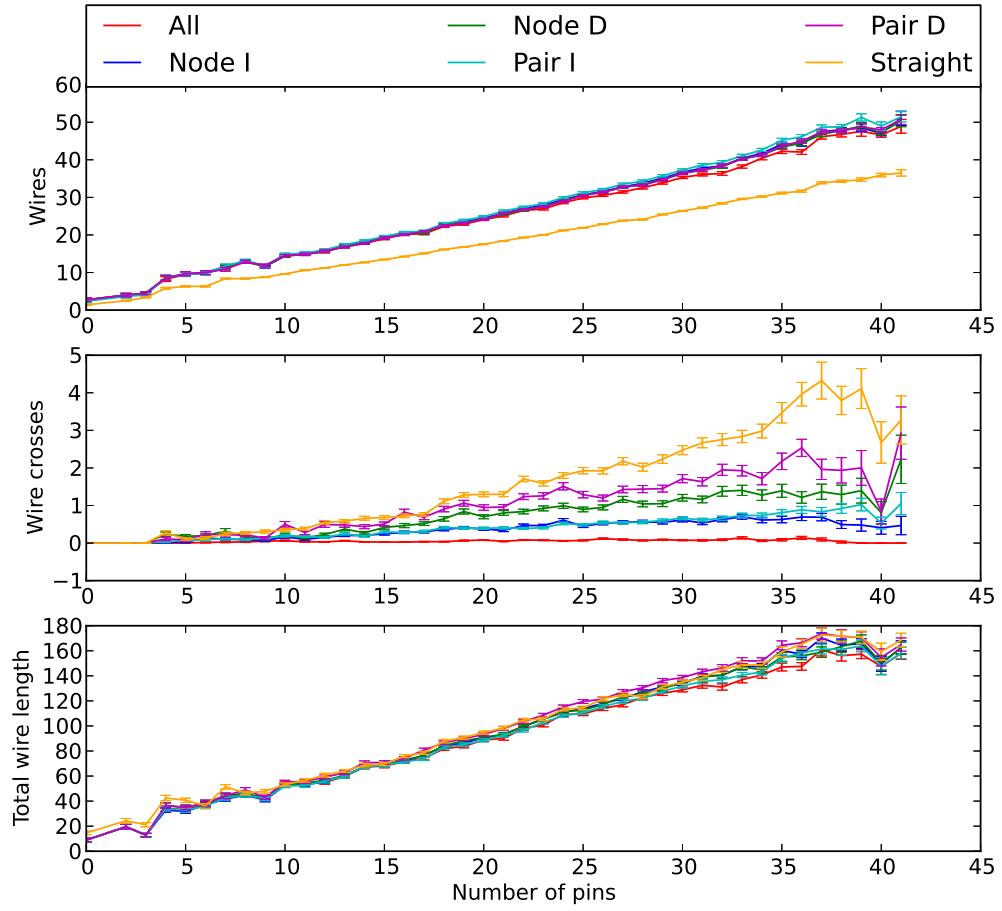


Figure 4-20: Wiring method comparison: layout quality trends.

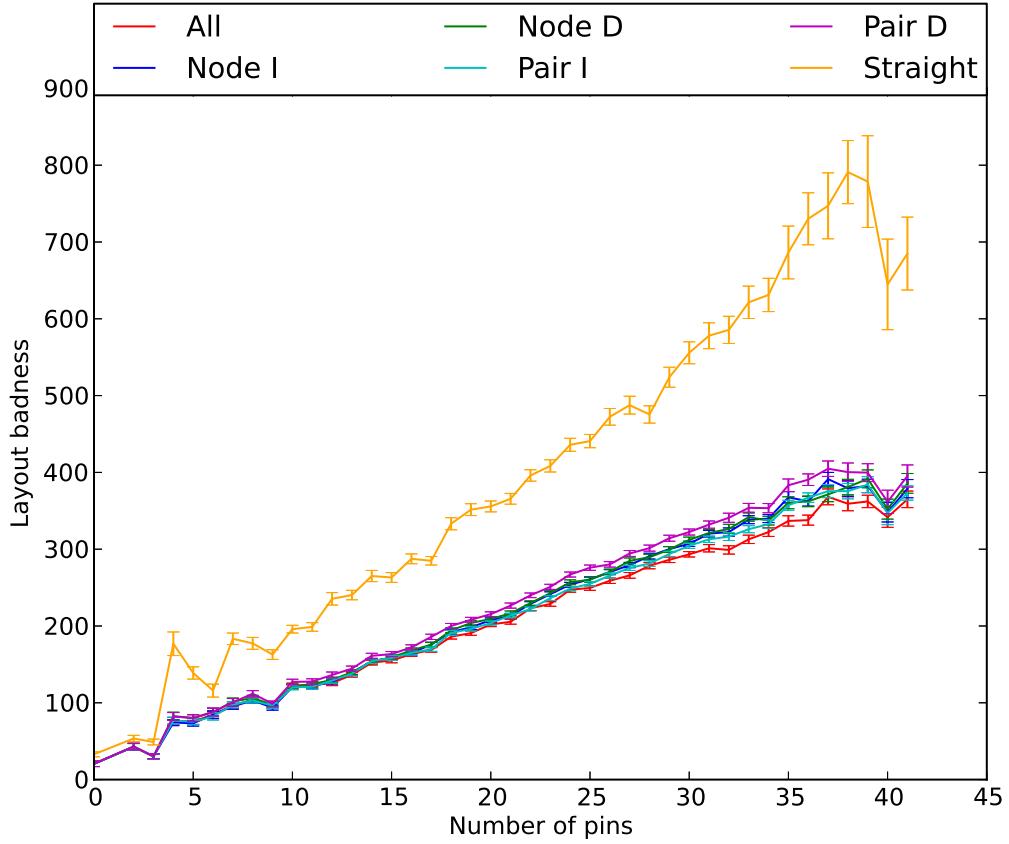


Figure 4-21: Wiring method comparison: layout badness trends.

4.3 Effect of limiting the number of vertices to expand in A^*

For each of the 5 wiring methods we have looked at, we have run tests on 4425 schematics. We have tested each method 10 times on each schematic. Hence, we have 44250 data points for each of the wiring methods. Here we provide histograms that show the maximum number of expanded vertices in the A^* search among the 44250 runs for each of the 5 methods.

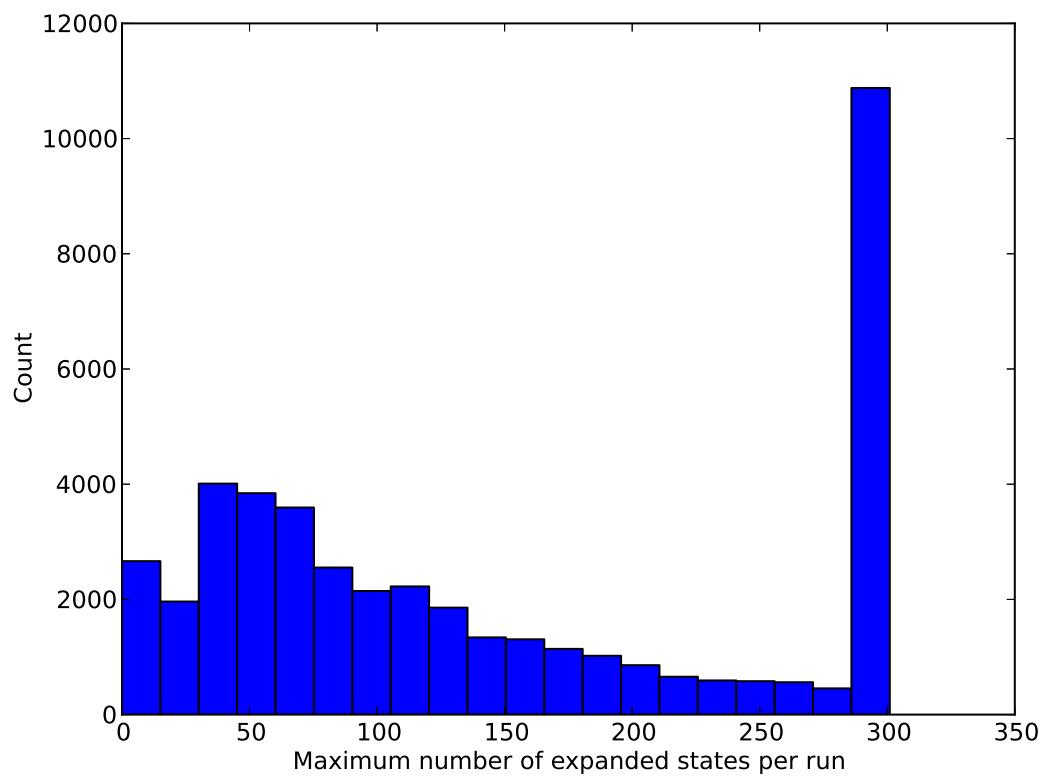


Figure 4-22: Histogram of number of vertices expanded for all pairs wiring.

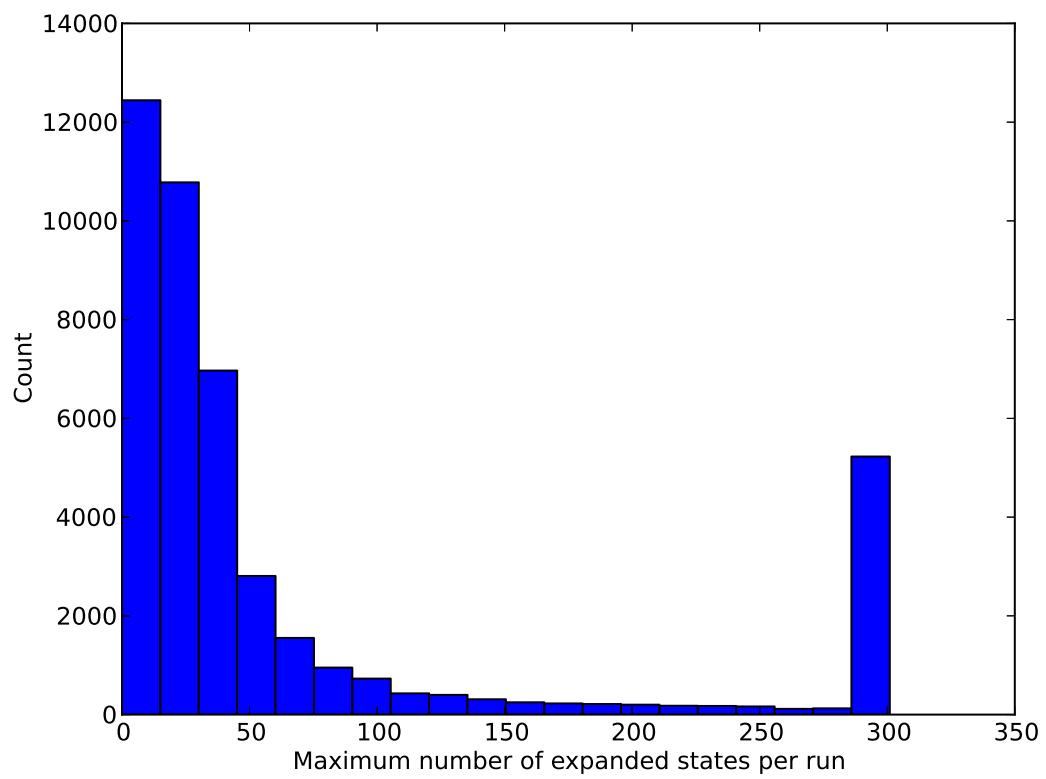


Figure 4-23: Histogram of maximum number of vertices expanded for per-node (increasing) wiring.

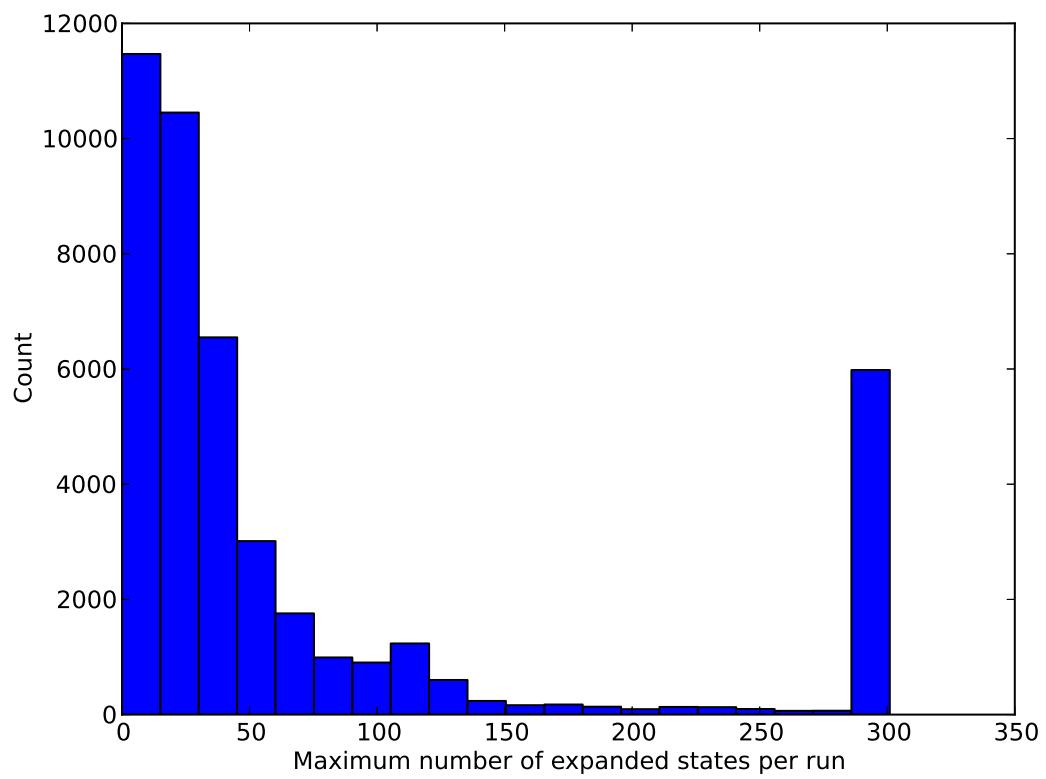


Figure 4-24: Histogram of maximum number of vertices expanded for per-node (decreasing) wiring.

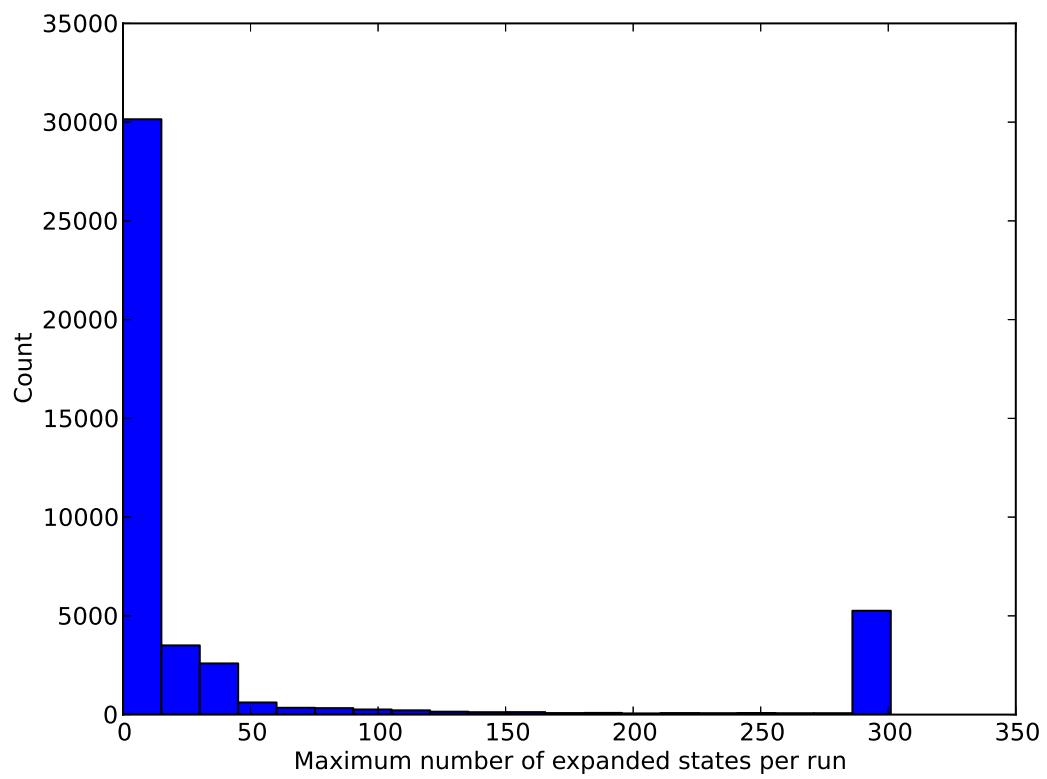


Figure 4-25: Histogram of maximum number of vertices expanded for per-pair (increasing) wiring.

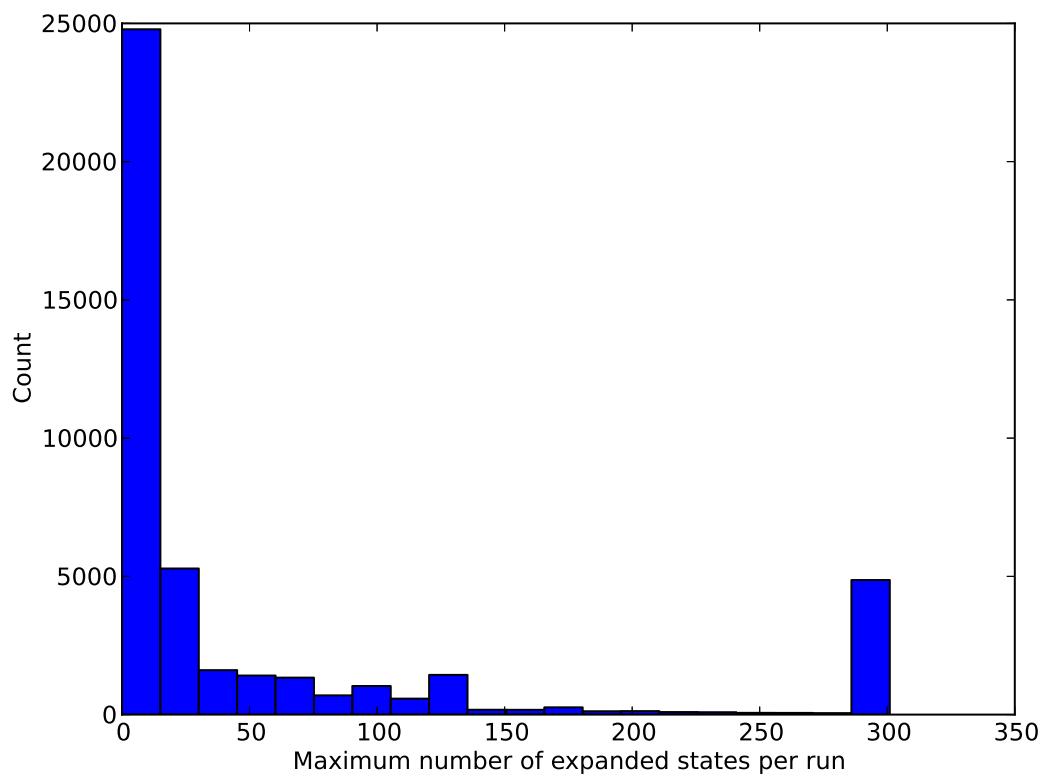


Figure 4-26: Histogram of maximum number of vertices expanded for per-pair (decreasing) wiring.

4.4 Comparing Search Methods

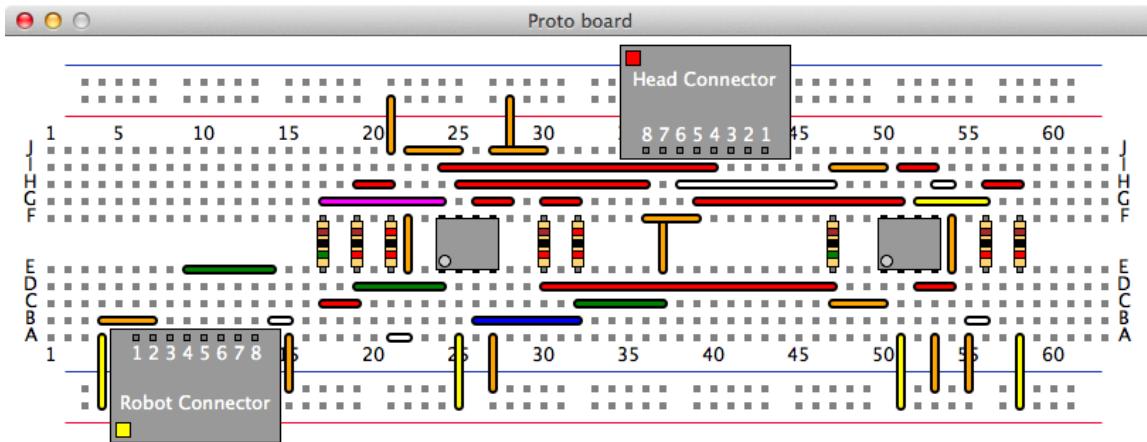


Figure 4-27: A^* Search exemplar.

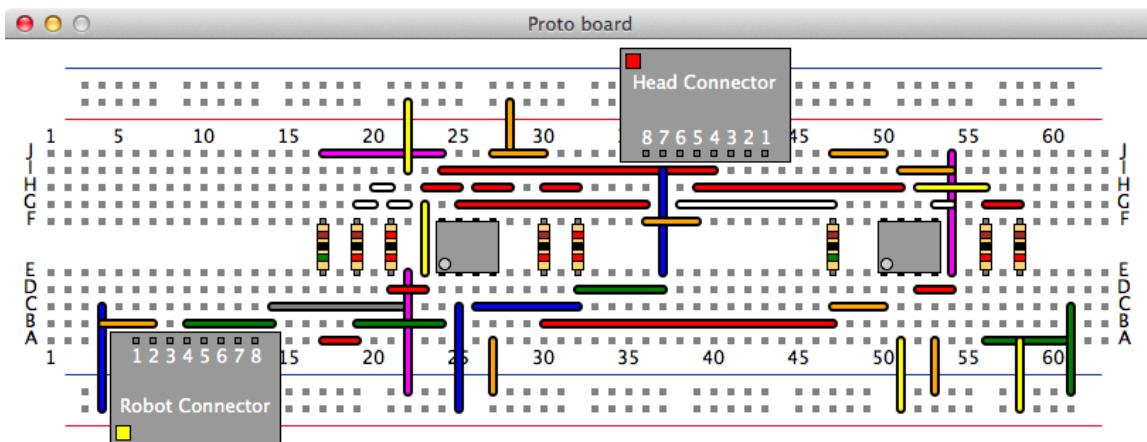


Figure 4-28: Best First Search exemplar.

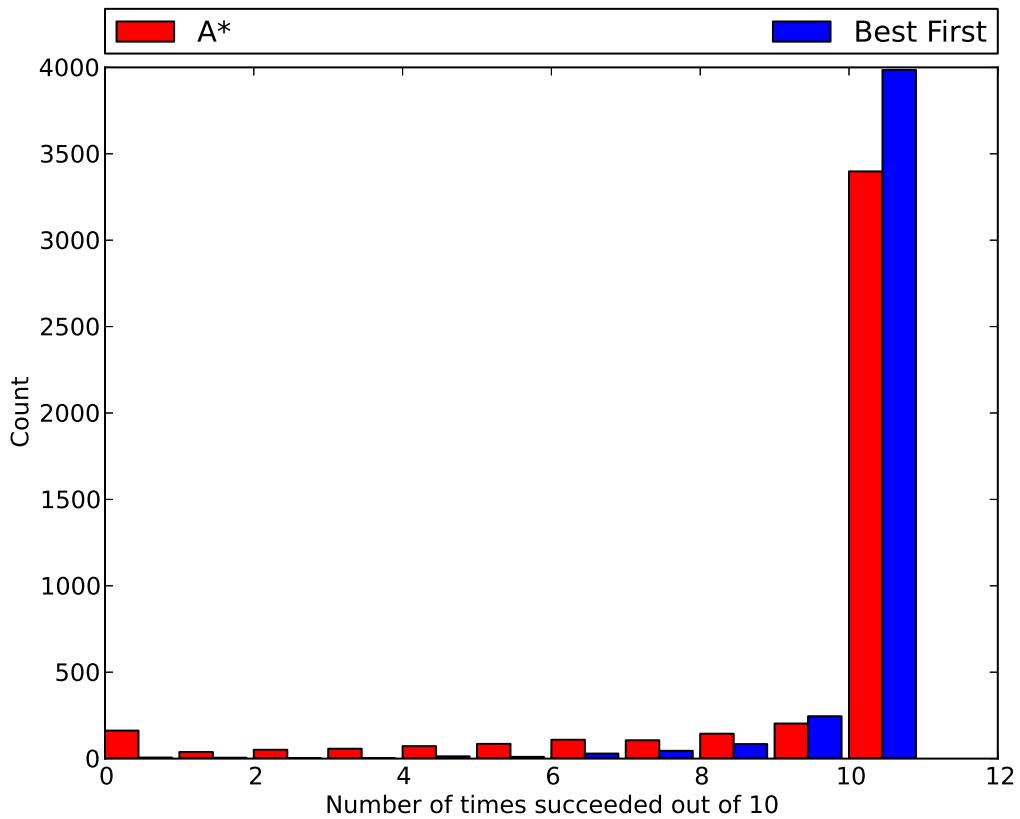


Figure 4-29: Search method comparison: success rates.

		Number of times succeeded out of 10										
		0	1	2	3	4	5	6	7	8	9	10
<i>A*</i>	Count	162	38	51	57	72	85	109	106	144	203	3398
	Probability	0.04	0.01	0.01	0.01	0.02	0.02	0.02	0.02	0.03	0.05	0.77
Best First	Count	6	5	2	1	13	10	29	45	84	245	3985
	Probability	0.00	0.00	0.00	0.00	0.00	0.00	0.01	0.01	0.02	0.06	0.90

Table 4.3: Search method comparison: success rates.




Figure 4-30: Search method comparison: success rate trends.

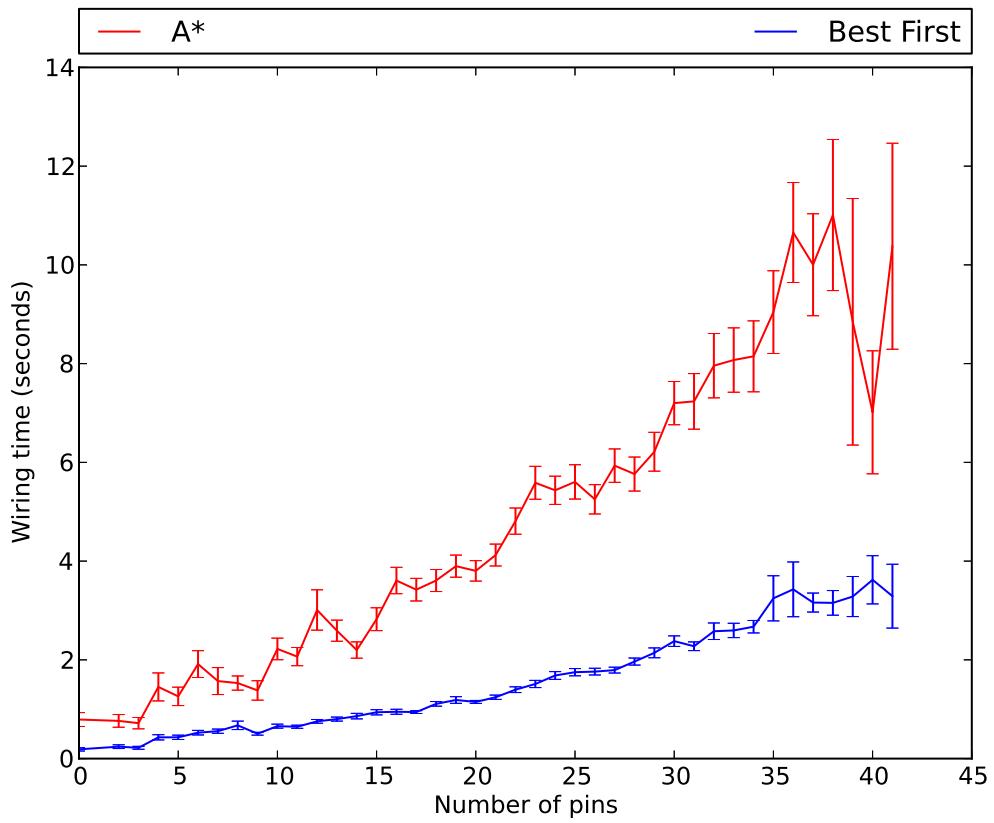


Figure 4-31: Search method comparison: wiring time trends.

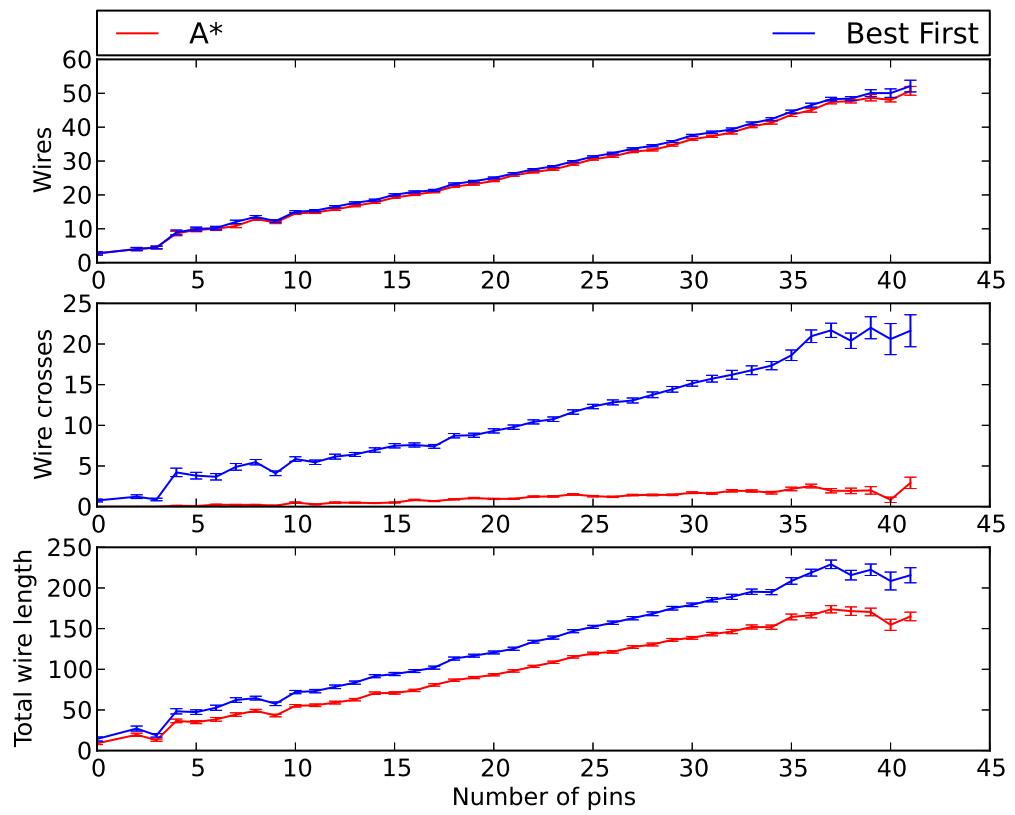


Figure 4-32: Search method comparison: layout quality trends.

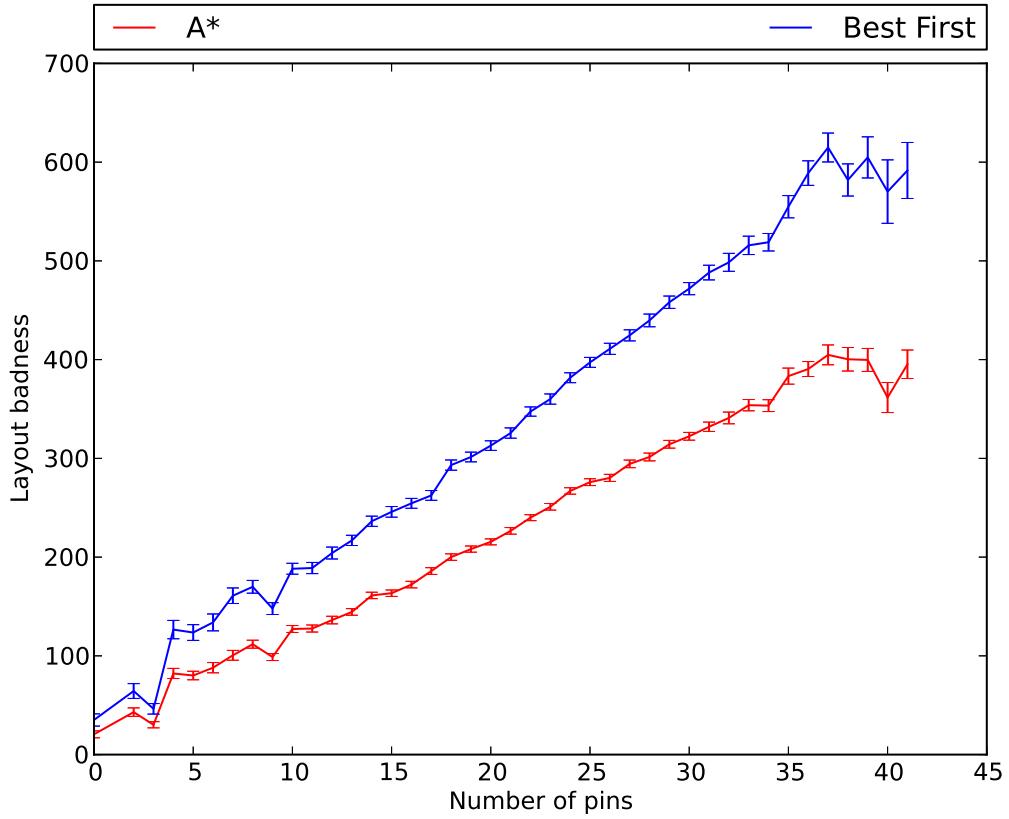


Figure 4-33: Search method comparison: layout badness trends.

4.5 Combined Algorithm

Here we provide data for the combined algorithm presented in Section 3.2.3. To generate this data, we used a different dataset of 4425 schematics. As desired, the algorithm has a 100% success rate. Figure 4-35 gives a breakdown of how the algorithm succeeded. The first four columns correspond to success from one of the four combinations of placement and wiring methods. The last 5 columns correspond to layouts in which none of the four combinations was successful on all pairs of locations and we had to connect a few pairs of locations by putting down a straight wire bridging the locations. Figure 4-36 gives the average total time taken by the algorithm as a function of circuit complexity. Finally, Figures 4-37 and 4-38 give statistics on the

quality of the layouts produced by the combined algorithm as a function of circuit complexity.

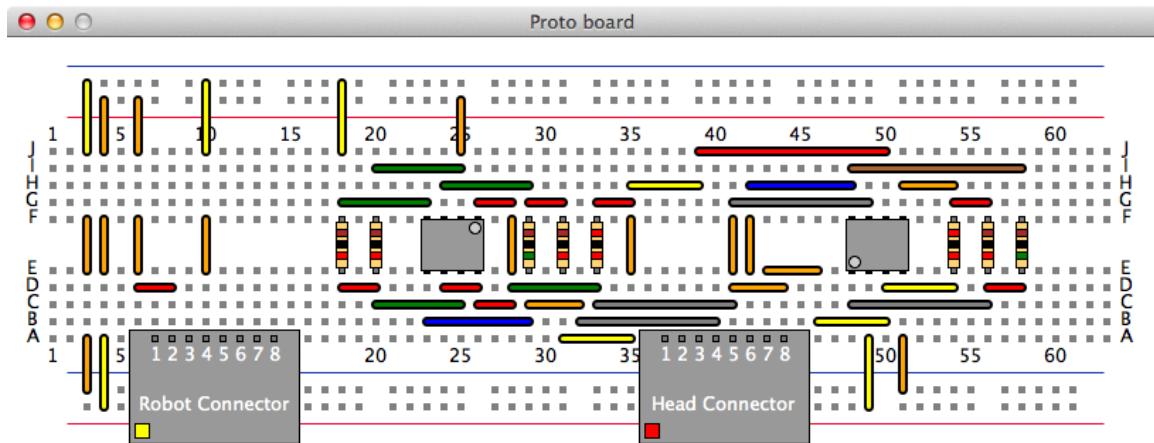


Figure 4-34: Combined algorithm exemplar.

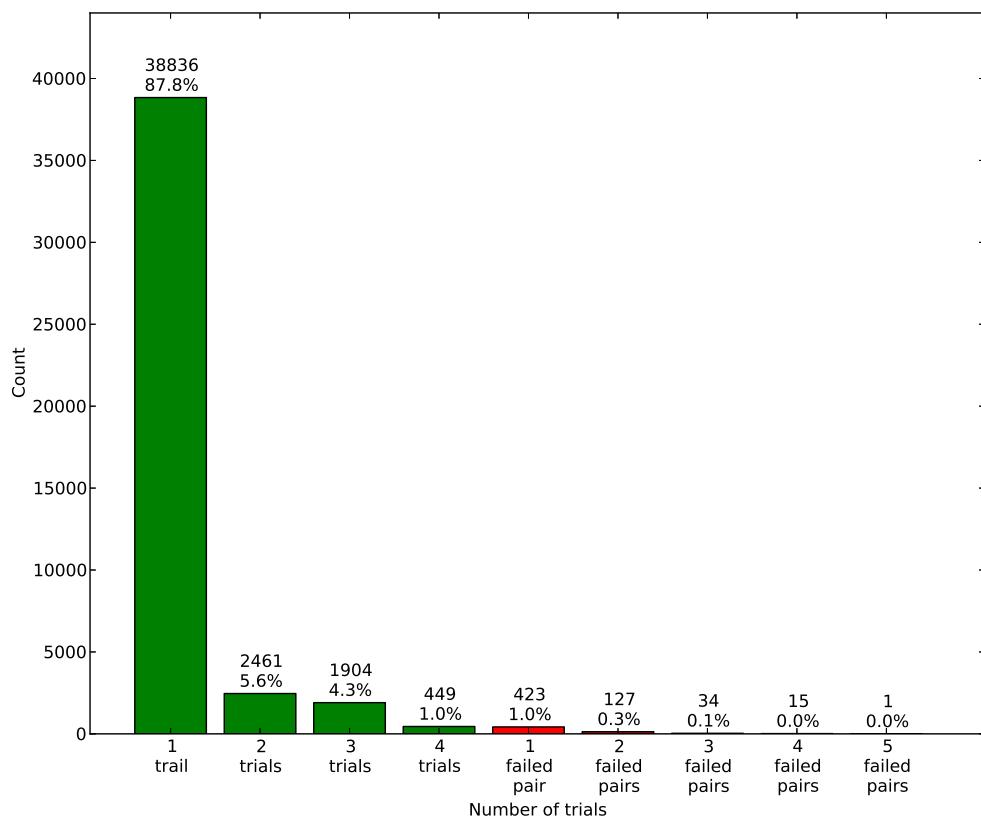


Figure 4-35: Combined algorithm success summary.




Figure 4-36: Combined algorithm total CPU time trend as a function of circuit complexity.

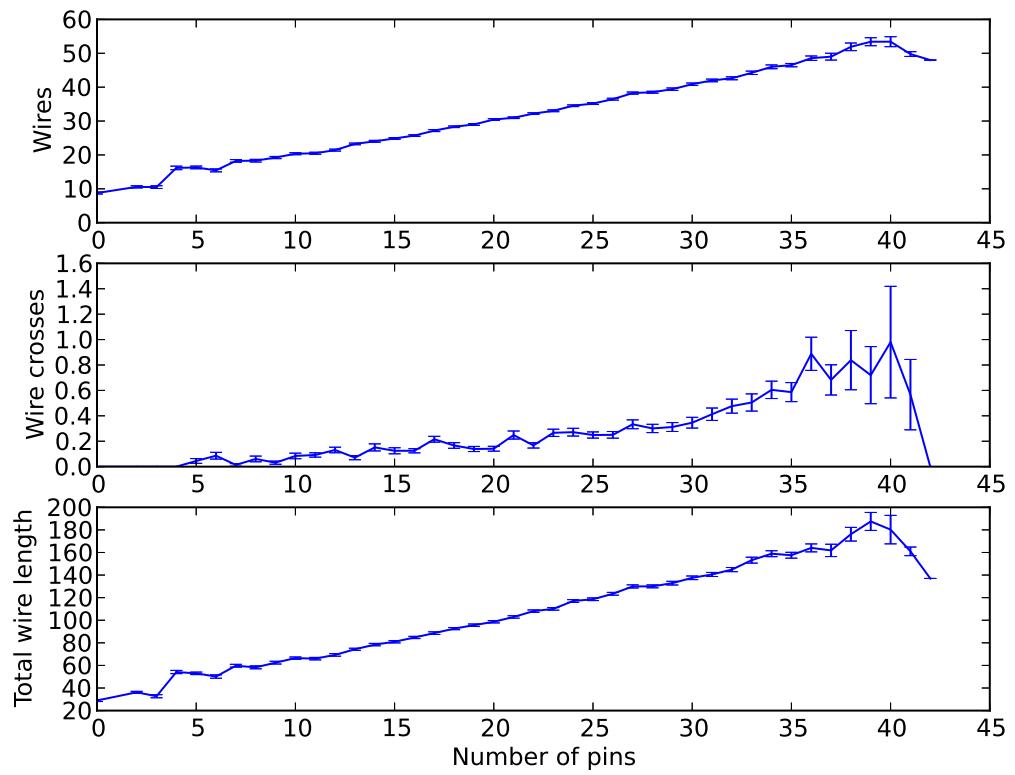


Figure 4-37: Combined algorithm quality trend.

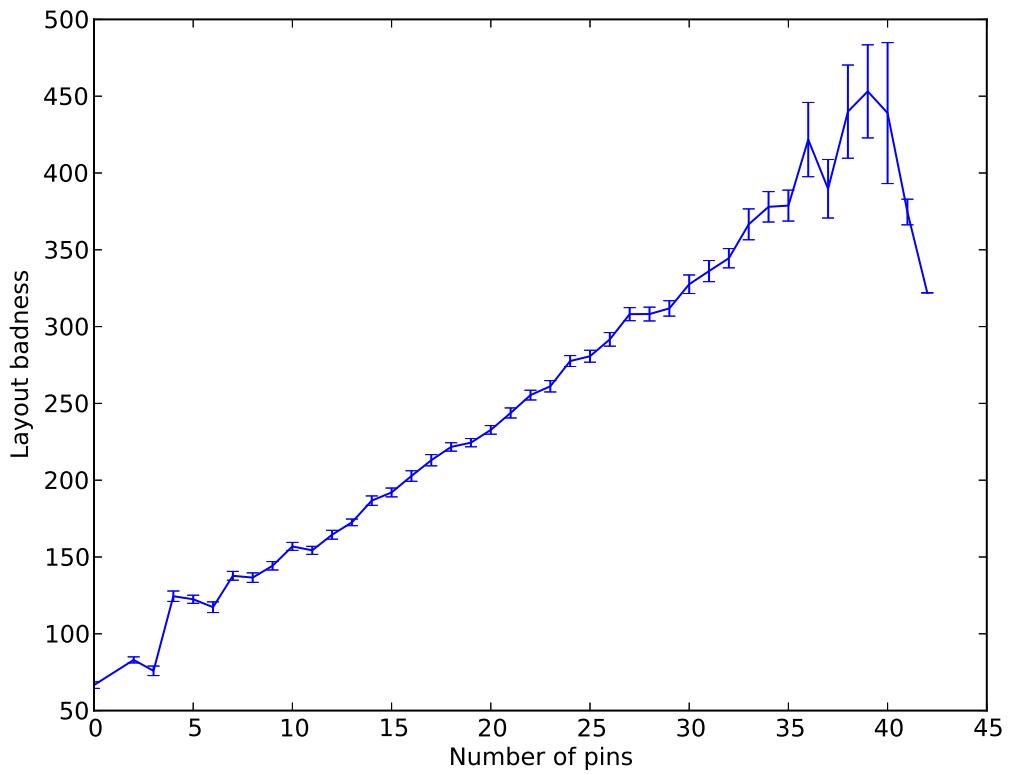


Figure 4-38: Combined algorithm layout badness trend.

Chapter 5

Discussion

In this chapter we provide justifications for the choices made in solving the protoboard layout problem, as well as detailed analysis of the data presented in Chapter 4.

5.1 Search Space Size

The proposed solution to this problem involves several simplifications and uses of heuristics. This is a result of the fact that the search space we are working with is very large. It is difficult to say exactly how large this search space is, but we can get an idea of its size. Let us just consider the number of ways we can put wires down on an empty protoboard (even in ways that may not make sense from a circuit theory standpoint). Finding this number reduces to finding the number of ways $T(n)$ in which we can choose pairs out of n items. Equation 5.1 gives an expression for $T(n)$ ¹.

$$T(n) = \sum_{k=0}^{\lfloor \frac{n}{2} \rfloor} \frac{n!}{k!(n-2k)!2^k} \quad (5.1)$$

In this problem, we have that $n = 830$, the number of available locations on an empty protoboard. Evaluating T at $n = 830$, we get approximately $2.8e1043$. The

¹The sequence of numbers described by $T(n)$ is sometimes referred to as the telephone numbers or the involution numbers.

largeness of this number indicates that doing any sort of exhaustive search will be hopeless.

5.2 Justifying Placement Choices

Resistors

For the sake of simplicity, and to significantly reduce the search space, we place resistors in the middle strip of the protoboard as shown in Figure 3-2. This choice is critical, as the resistor pieces can generally be placed at numerous places on the protoboard. With this restriction, there are 63 slots available for one resistor. Without this restriction, there are a total of 763 slots available. The restriction is good when we consider the reduction in the search space size. On the other hand, this restriction is bad as it imposes a restriction on the size of the schematics (in terms of the number of components) for which a layout can be generated using the algorithm. Given that the number of resistors in a typical 6.01 circuit is very small, this restriction proves to be very useful.

Op-amps

Op-amps are the trickiest components to handle because each op-amp package put on the protoboard contains two op-amps within it. Equation 5.2 presents an expression for the number of different ways to package together n op-amps. For example, if we have 2 op-amps, we can either use one op-amp package for each, or put them both in the same package, which we can do in one of two different ways. All together, there are 3 different ways to package together 2 op-amps. Table 5.2 gives the number of different packagings possible for various n .

$$\text{Number of ways to package } n \text{ op-amps} = \sum_{k=0}^{\lfloor \frac{n}{2} \rfloor} \frac{n!}{k!(n-2k)!} \quad (5.2)$$

Our placement approach explores all possible ways of packaging the op-amps. We do this because the typical 6.01 circuit contains no more than 6 op-amps, and so we

n	Number of ways to package n op-amps
1	1
2	3
3	7
4	25
5	81
6	331
7	1303
8	5937
9	26785
10	133651

Table 5.1: Number of ways of packaging together n op-amps for various values of n .

are tasked with exploring at most 331 alternatives, which is not too computationally intensive. If the algorithm was meant to handle a larger number of op-amps (for example 10), then this approach would not work as the number of alternatives to consider would be far too large, and we would have to consider different strategies.

5.3 Explaining the Results

Chapter 4 presented quantitative data to compare alternative strategies for solving the protoboard layout problem. Here, we analyze those data and give reasonings for why we obtained those results.

5.3.1 Comparing Placement Methods

Success Rate

The blocking placement method (89.0% overall success rate) is slightly more successful than the distance placement method (87.6% overall success rate). While these alternative methods are not significantly different in terms of success rate, we note that both methods are significantly more successful than the random placement alternative (43.5% overall success rate). As a function of circuit complexity, Figure 4-7 suggests that the two alternatives have almost identical success rates. As we would

expect, success rate generally decreases for both of the placement methods as circuit complexity increases. Once again, we observe that the distance and blocking based placement methods are much more successful than the random placement method, especially as circuit complexity increases. The rise in success rate for high circuit complexity is a result of the fact that there are very few circuits of the highest complexity, on which the algorithm happened to be consistently successful.

Wiring Time

We observe from Figure 4-8 that, once again, the two methods are very similar, with random placement being markedly worse than both. We see that the distance method generally takes slightly less time than the blocking method, but the difference between the two is almost negligible. As we would expect, we see that as the complexity of the circuits increases, the amount of time spent by the wiring step also increases. As with the success rate trends as a function of circuit complexity, we observe that there are outliers at the far end of the figure due to a very small sample of the most complex circuits in the randomly generated schematic dataset. As circuit complexity increases, the variance in wiring time when using random placement also increases significantly. This is a result of the fact that random placement frequently fails on the most complex circuits, and so the sample size we have for successful runs is very small.

Layout Quality

Figure 4-9 presents graphs that compare numbers of wires, numbers of wire crosses, and total wire lengths. Figure 4-10 shows the trend of layout badness computed using our metric as a function of circuit complexity. We first observe that the number of wires used by the two methods are almost identical. When we look at the number of wire crosses in the layouts, we see that the blocking method consistently results in more wire crosses. Similarly, when we look at the total length of wires used in the layouts, the blocking method exceeds the distance method consistently, with the difference getting higher as circuit complexity increases. This is somewhat expected

because the blocking method does not directly aim to put pieces that need to be connected close together whereas the distance method directly tries to minimize the total wire length that may be needed. In terms of the layout badness metric, we observe once again that the distance method is slightly better than the placement method for almost all circuit complexity values. Finally, we observe that the random placement method, per our badness metric, produces much worse layouts than the other two methods when it does succeed.

Conclusion

It is difficult to conclusively pick the best placement method from these results. What we can determine is that the distance based method and the blocking based method are both significantly better than random placement.

5.3.2 Comparing Wiring Methods

Success Rate

Figures 4-17 and 4-18, and Table 4.2 show that the all pairs wiring method has a much smaller success rate than all of the other alternatives, especially as circuit complexity increases. The reason why the all pairs method has a smaller success rate is in large part due to the maximum allowed number of vertices to expand in A^* . The figures in Section 4.3 depict that this limit affects the all pairs wiring method much more than it does the others. This result is expected because the search task in all pairs wiring is more difficult than in the other alternatives. Notice that as the search tasks get less difficult from all pairs to per-node to per-pair wiring, the effect of the expanding limit decreases. The other four alternatives have very comparable success rates.

Wiring Time

Figure 4-19 compares wiring time across the five methods. Once again, we observe that the all pairs methods takes significantly more time than the other methods. We also observe from Figure 4-19 that the wiring times for the two per-node methods

are comparable, and that the wiring times for the two per-pair methods are also comparable, but that the per-node wiring times are generally bigger than the per-pair wiring times. This trend is also expected as the per-node methods attempt to connect multiple pairs of locations at once, which generally requires searching through more alternatives than connecting each of the pairs of locations individually.

Layout Quality

Figures 4-20 and 4-21 compare the quality of the layouts produced by the five alternative wiring methods. First, as we observed in the placement method comparison, we see that there is very little difference in terms of number of wires used and the total wire length. However, there are noticeable differences in the number of wire crosses. We see that the all pairs method generates layouts with much fewer wire crosses than the other methods. This is expected since the algorithm runs one search to connect all pairs of locations. Conversely, the per-pair (decreasing) and per-node (decreasing) methods result in the largest number of wire crosses. Note that the per-node (decreasing) method produces more wire crosses on average than the per-pair (increasing) method. We observe that the order in which we consider pairs of locations has a telling effect on how good the generated layouts will be. In essence, connecting the harder pairs of locations generally produces more wire crosses.

Conclusion

While the all pairs method is the least successful method and generally takes the longest among the five, it tends to produce the best layouts when it does succeed. On the other hand, the alternatives that break the problem down into smaller pieces succeed more often and finish more quickly, though they tend to produce worse results. Furthermore, the more finely we break down the problem, the faster the overall algorithm runs. Lastly, ordering subproblems from hardest to easiest has the effect of making the overall wiring step faster, but produces worse results than the reverse order and does not result in a markedly better success rate.

5.3.3 Comparing Search Methods

Success Rate

Figures 4-29 and 4-30 and Table 4.4 present data comparing the success rates of the two search algorithms. We observe that Best First Search is much more successful than A^* . 98% of the test circuits were solved at least 8 times out of 10 when we used Best First Search, versus 85% when we used A^* . This result is not surprising because when using Best First Search the algorithm looks for layouts that satisfy the connection requirements, ignoring the badness of the layouts it considers. Hence, Best First Search is less susceptible to the restriction on number of vertices to expand than is A^* .

Wiring Time

The fact that Best First Search settles for any layout that satisfies the connection requirements suggests that it should finish more quickly in addition to being more successful. Figure 4-31 supports this expectation.

Layout Quality

Figure 4-32 shows that the layouts generated by the algorithm when using Best First Search are worse than the layouts generated when using A^* . Most importantly, the number of wire crosses in the layouts produced by Best First Search are markedly greater than the number of wire crosses in the layouts produced by A^* . We also observe that the total wire length is greater when using Best First Search.

Conclusion

Our choice of a search algorithm forces us to consider a trade-off between speed, success rate, and quality. Using Best First Search, most runs will be successful and terminate quickly, but will produce very poor results. Using A^* , fewer runs will be successful, and the successful runs will take longer to terminate, but the resulting layouts will generally be better.

5.3.4 Putting Them All Together

In this section we discuss why Algorithm 2 is structured the way it is, and we also discuss the data we obtained for the combined algorithm. Recall that the combined algorithm makes four attempts at generating a complete layout. We try both placement methods, using the distance method first and the blocking method second. We try the distance method first because it tends to generate layouts with fewer crossing wires and smaller total wire length. We use per-pair wiring, and consider both orders of doing the wiring, increasing order first and decreasing order second. We use per-pair wiring because it takes considerably less time than both per-node and all pairs wiring, and neither of the other wiring methods has a better success rate. We try increasing order first because it tends to generate layouts with fewer crossing wires then decreasing order.

Let us now consider the data we obtained for the combined algorithm. Firstly, we saw that the combined algorithm had a 100% success rate, which is critical when we consider the fact that students will be using a tool that almost never fails. This success rate is a result of the last part of Algorithm 2 that uses single wires to connect pairs of locations that need to be connected. Figure 4-35 shows that, on the test dataset of 4425 schematics, with 10 runs carried out on each schematic, the algorithm succeeded in generating layouts for 98.6% of the circuits without having to put down forced wires (wires added at the last step of the algorithm to connect any disconnected location pairs). The algorithm was required to put down more than 2 forced wires on only 0.1% of the circuits. From Figure 4-36, we see that as the circuit gets more complex, the amount of time the algorithm takes sharply increases. This is due to a more difficult placement task (which is most notable as the number of op-amps increases) as well as a more difficult wiring task. Importantly, the maximum point on the plot occurs at less than 180 seconds. This indicates that on the test dataset, the algorithm took at most about 3 minutes to run. This is encouraging from a practical standpoint because 3 minutes is not long for a student to have to wait for a layout to be generated. Finally, Figures 4-37 and 4-38 present trends of

quality as a function of circuit complexity. Most importantly, we observe that the average number of wire crosses on the test layouts never exceeded 1.

5.4 Further Work

5.4.1 Treating Resistors as Wires

Our current solution treats resistors as circuit components and so they are placed before the wiring search is executed. However, resistors have the special property among the components that they can be placed as if they were wires of a fixed set of lengths. This suggests we can think of resistors as wires, and thus can handle them in the wiring step of the algorithm instead of the placement step. As a result, the search that we do in the wiring step would now need to be more elaborate. Not only would we need to keep track of pairs of locations on the protoboard that need to be connected, but we would also need to know whether to put a resistor between the two locations. In the latter case, we have the restriction that one of the wires we use to connect the pair of locations needs to be of a length that can fit a resistor.

While this idea appears very promising, its implementation is not trivial. First, if there is any node in our circuit that is only connected to resistors (and no other circuit components), then that node will be unrepresented on the protoboard at the end of the placement step. One possible solution to this problem is to reserve an empty column on the protoboard for the node between the resistors, which we can then use in the wiring step. A better solution may be discovering the best places for the node as we are placing down resistors in the wiring step, but this solution would make the search we carry out more complicated. Second, we will need a new kind of heuristic for our *A** search that takes resistors into account. One possible solution is to highly penalize an absent resistor when computing the heuristic. Even with an amended heuristic, however, there are cases where the search may not find an answer (for instance a pair of adjacent locations that must be connected by a resistor).

5.4.2 Building Layouts Similar to Previously Generated Layouts

One of the weaknesses of our tool is that it may generate different layouts for the same circuit schematic. This is not a result of an explicit randomization in our algorithm, but rather a side effect of some data structures we use such as Python sets and dictionaries. A similar problem is that a slight change in the circuit schematic may result in a completely different layout. It is, therefore, very important that students in 6.01 be confident that they designed the right circuit before taking the time to build their circuit. This is where the simulation capabilities of our tool come in to play. As an alternative solution to this problem, it may be very useful to have our tool remember the last placement it used and try to produce a new placement as similar to the last one as possible.

5.5 Remarks

The results we have gotten for the final algorithm are encouraging. It has a 100% success rate and produces layouts that, by our metric, are very good. If this tool saves 1 lab hour per student per semester (time the student would have spent producing layouts), then we are looking at a savings of about 200 6.01 lab hours per semester, which is fantastic success.

Appendix A

Schematic Drawing GUI

One of the aims of this project is making it easy for students to describe the circuit they have in mind to a computer. To that end, it is important that the schematic entry tool be intuitive and easy to use. In this section, we describe the capabilities and features offered by the schematic entry GUI. Figure A-1 presents the GUI containing a sample schematic. The figure depicts the four important sections of the GUI: the board, the palette, the analysis section, and the cursor toggle section.

A.1 Palette

The palette (item 3 in Figure A-1) offers all of the circuit components that can be included in a 6.01 circuit. Clicking a circuit component on the palette spawns a new component of the same type on the board above the palette (the board is described in detail in Section A.2). This component can then be used in the circuit construction. The “Robot” and “Head” buttons on the palette spawn multiple parts at once corresponding to the multiple parts contained within the robot and head connectors, respectively. The robot connector contains pins for power and ground as well as for analog input pins and one analog output pin. The head is composed of a motor, a motor pot, and two photosensors. Figure A-2 shows these last two circuit components as they would appear in the editor.

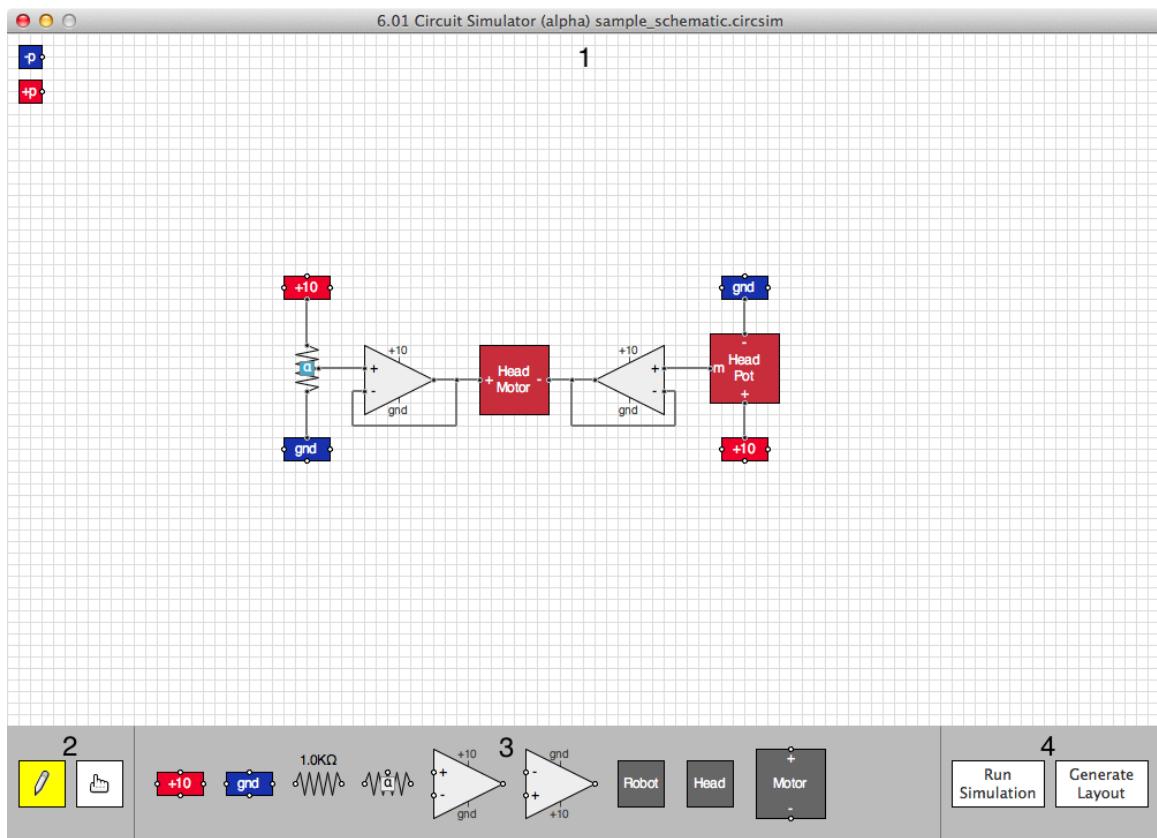


Figure A-1: Parts of the schematic entry GUI.




Figure A-2: Grouped components: Robot Connector and Head Connector.

A.2 Board

The board (item 1 in Figure A-1) is where the user can draw circuit schematics. The user may move a component on the board by clicking the component and dragging it to the desired place. When dragging, the GUI draws guide lines that extend to the edges of the board to help the user place the component at the right place. The user has the option to select and move multiple items at once. The user may delete a component by clicking on it while pressing Ctrl. The user may rotate a component by clicking on it while pressing Shift. An important aspect of circuit schematics is interconnecting components with wires. Each circuit component in the GUI comes with a few connection points. The user may draw a wire to connect components by clicking on a connection point on a component and dragging. A wire may be drawn to another connection point, or a wire already on the board (which snaps the new wire onto the existing wire), or an empty location on the board (which would create a new connection point). A wire may also be drawn starting from an existing wire, which creates a new connection point on the existing wire. The GUI allows the user to drag connection points. To achieve this, the GUI has two possible states for the cursor, the drawing state and the dragging state (mainly referring to wires). Item 2 in Figure A-1 displays the panel that lets the user toggle between these two states. In the dragging state, the user can drag connection points just like other circuit components. When drawing wires, the GUI attempts to route the wires in a way that is aesthetically pleasing. That is, the wire is routed so as to avoid crossing wires and, more importantly, wires crossing components on the board. This routing is not a trivial task. In fact, this problem is very similar to the layout problem that this project aims to tackle. The solution to the wiring problem in the GUI also uses search.

A.3 Analysis

The analysis section of the GUI (item 4 in Figure A-1) lets the user analyze the drawn circuit schematic in two ways.

A.3.1 Simulation

The GUI lets the user simulate the circuit and test whether it behaves as expected. The simulation infrastructure is ported from CMax, so circuits are simulated exactly as they are simulated by CMax. If there are probes in the circuit, the simulator presents the voltage difference across the probes as an output. If there is a motor in the circuit, the simulator presents the motor's angle and motor's rotational velocity as functions of time. If there are any pots in the circuit, the user is expected to select a simulation file for each pot describing how the pot is manipulated as a function of time. Similarly, if the photosensors are a part of the circuit, the user is expected to select a simulation file for each photosensor set describing the corresponding lamp's distance and angle from the head. These simulations help students (and staff) verify that they have a correctly functioning circuit before building it.

A.3.2 Layout

The GUI also lets the user generate a layout for the circuit schematic, which is the main object of this project. Very importantly, the GUI makes it very easy to relate the schematic with the layout. When the user hovers over a component in either window, the GUI highlights the corresponding component in the other window. When the user hovers over a wire in one window, the GUI highlights all of the wires in both windows that correspond to the same node. Figures A-3 and A-4 demonstrate this feature.




Figure A-3: Wire highlighting example.




Figure A-4: Component highlighting example.

A.4 Other Features

Here we discuss several features offered by the GUI that have not been discussed so far:

1. The GUI allows the user to save schematics for later viewing or editing.
2. Protoboard layouts can also be saved as CMax files allowing for editing in CMax.
3. The schematic editing tool allows the user to undo and redo all actions.
4. The GUI has menu items that offer access to some of the features already discussed. Menu items are added when particular circuit components are selected. For instance, selecting a pot component results in a new menu item that allows the user to select a signal file for the pot. The same properties can also be reached by right-clicking on the components.
5. The GUI changes the cursor appropriately to provide feedback. For instance, the cursor becomes a pencil if the user can draw a wire starting at the cursor's current position. The cursor also changes to indicate when the user is about to rotate or delete a component. If the tool is busy either running a simulation or generating a layout, the cursor shows that the tool is busy.

A.5 Shortcuts

Table A.5 presents the shortcuts available in the GUI.

Action	Shortcut
Ctrl + n	New file
Ctrl + o	Open file
Ctrl + s	Save file
Ctrl + q	Quit
Ctrl + z	Undo
Ctrl + y	Redo
Ctrl + w	Close simulation windows
g	Generate layout
s	Run simulation
Delete	Delete selected items
r	Rotate selected item
d	Toggle cursor state
← h	Move selected items left
↓ j	Move selected items down
↑ k	Move selected items up
→ l	Move selected items right

Table A.1: Shortcuts.

Bibliography

- [1] Introduction to EECS I. <http://web.mit.edu/6.01>.
- [2] 6.01 Staff. Cmax.